

Act or Escalate? Evaluating Escalation Behavior in Automation with Language Models

Anonymous authors

Paper under double-blind review

Abstract

Effective automation hinges on deciding when to act and when to escalate. We model this as a decision under uncertainty: an LLM forms a prediction, estimates its probability of being correct, and compares the expected costs of acting and escalating. Using this framework across five domains of recorded human decisions—demand forecasting, content recommendation, content moderation, loan approval, and autonomous driving—and across multiple model families, we find marked differences in the implicit thresholds models use to trade off these costs. These thresholds vary substantially and are not predicted by architecture or scale, while self-estimates range from well-calibrated to systematically miscalibrated in opposite directions across models. We then test interventions that target this decision process by varying cost ratios, providing accuracy signals, and training models to follow the desired escalation rule. Prompting helps mainly for reasoning models. SFT on chain-of-thought targets yields the most robust policies, which generalize across datasets, cost ratios, prompt framings, and held-out domains. These results suggest that escalation behavior is a model-specific property that should be characterized before deployment, and that robust alignment benefits from training models to reason explicitly about uncertainty and decision costs.

1 Introduction

Agents based on large language models (LLMs) are increasingly deployed to automate consequential decisions, from performing autonomous tasks (Wang et al., 2024) to generating code (Jimenez et al., 2024) and managing enterprise workflows (Wu et al., 2024). Most evaluations focus on speed, accuracy, and cost-savings (Jimenez et al., 2024; Chiang et al., 2024; Trivedi et al., 2024). However, in each setting, an agent faces a choice that receives less direct evaluation: should it implement its own decision, or escalate to a human? Existing LLM evaluations measure what models *can do*; we characterize what they *choose to do* under uncertainty. This act-or-defer policy is not captured by accuracy alone, and it can leave a high-accuracy model either propagating errors at scale or eliminating the value of automation.

Agent escalation behavior is fundamental to effective automation. An agent that fails to escalate when it is uncertain or incorrect will propagate errors at scale, while an agent that always escalates provides no labor savings over manual review. Two parameters govern effective escalation. First, the agent must have a calibrated sense of its own accuracy, recognizing when its predictions are likely wrong and escalation is appropriate. Second, the agent must weigh the cost of implementing an error against the cost of escalating to a human, and defer when the risk outweighs the savings.

In the main text we report results for eight models in four families, each represented by a smaller and larger variant: Qwen3.5-4B and Qwen3.5-9B, Gemma 3 4B and Gemma 3 12B, Claude Sonnet 4.6 and Claude Opus 4.7, and GPT-5-nano and GPT-5-mini. The Appendix extends the analysis to five additional models: the larger MoE flagship Qwen3.5-397B-A17B, plus Llama 4 Maverick and Llama 3.3 70B, and Mixtral 8x7B and Mistral Small 24B. We evaluate these models on five decision-making tasks, each derived from large-scale




Figure 1: The escalation decision. A human delegates a task to an LLM agent, which produces a prediction and then decides whether to implement (incurring error cost c_w if wrong) or escalate (incurring labor cost c_ℓ). The agent escalates when its estimated probability of being correct $\hat{p}$ falls below a threshold τ .

human decision data: demand forecasting, loan approval, content moderation, content recommendation, and moral dilemmas (featured in the Appendix). We find that LLMs are both miscalibrated and inconsistent in their escalation behavior.

First, models differ widely in self-assessment of their own predictive accuracy. Some systematically overestimate it, others underestimate it, and others are well-calibrated. The direction and magnitude of miscalibration vary by model and domain. Prior work has examined whether LLMs can assess their own knowledge (Kadavath et al., 2022) and documented calibration failures in question answering (Xiong et al., 2024), but these studies measure confidence in isolation. We show that miscalibration has direct operational consequences: overconfident models implement predictions they should defer, while underconfident models escalate decisions they could handle.

Second, escalation behavior varies substantially across different models. This variance is not associated with model architecture or size. Some models tend to prefer escalating (lower decision threshold), some generally favor implementing (higher decision threshold). This variation persists even within model families. Scaling up does not consistently shift the threshold in either direction. Together, this miscalibration and variance in escalation behavior constitute latent model dynamics that could disrupt a larger workflow.

These dynamics, however, can be corrected. We test a series of interventions that target the escalation decision. Prompt-based cost framing alone is insufficient, but combining it with extended thinking produces larger improvements for reasoning models. Supervised fine-tuning (SFT) on chain-of-thought targets proves most effective: the resulting models make high-accuracy escalation decisions across the evaluated datasets, cost ratios, and prompt framings, and generalize to held-out domains they were never trained on. Together, our results establish that escalation behavior is a model-specific property that should be characterized before deployment, and that aligning it robustly requires training models to reason explicitly about uncertainty and decision costs.

2 Theory

2.1 The escalation decision

Consider a workflow in which a human delegates a binary task to an LLM-based agent. The agent observes a scenario x and produces a prediction $\hat{y} \in \{0, 1\}$. The ground truth is the human’s preference $y \in \{0, 1\}$. The agent then decides whether to *implement* (act on) its prediction or *escalate* (defer to the human). We refer to this choice as the agent’s *escalation decision*: implementing commits to the agent’s own prediction, while escalating returns control to the human principal for that instance. Every action an agent takes in practice contains an implicit escalation decision: when it makes a tool call, generates a response, or commits a change, it is choosing to implement rather than escalate and ask for guidance. Figure 1 illustrates this workflow.

We model the agent’s behavior as follows. After producing its prediction, the agent forms an internal estimate $\hat{p}$ of its probability of being correct $P(\hat{y} = y)$. The cost structure then

83 determines an optimal decision threshold τ (derived below): the agent escalates if $\hat{p} < \tau$
 84 and implements otherwise.

85 2.2 Cost structure

86 One of two costs can be incurred depending on the agent’s action. If the agent escalates, the
 87 human pays a labor cost $c_\ell > 0$, which represents the time and effort required of the human.
 88 If the agent implements and its prediction is wrong, the human pays an error cost $c_w > c_\ell$.
 89 Correct implementations incur zero cost.

90 **Theorem 1** (Uniqueness of the optimal threshold). *Let the agent escalate whenever $\hat{p} < \tau$ for*
 91 *some threshold $\tau \in [0, 1]$, and suppose $\hat{p}$ is perfectly calibrated (i.e., $\hat{p} = p$). Then $\tau^* = 1 - c_\ell / c_w$*
 92 *is the unique threshold that minimizes expected cost.*

93 Proof in Appendix C. By straightforward extension, a calibration bias in which the agent
 94 perceives $\hat{p} = p + \mu$ shifts the effective threshold to $\tau^* - \mu$: overconfident models ($\mu > 0$)
 95 implement too aggressively and underconfident models ($\mu < 0$) escalate too often, and
 96 expected cost is increasing in $|\mu|$.

97 Theorem 1 establishes that practitioners cannot avoid characterizing their model’s threshold:
 98 any deviation from τ^* creates avoidable cost. Together with the calibration observation
 99 above, this motivates the empirical characterization we pursue in the following sections.

100 To first assess calibration, we give the agent no signal, so that its escalation rate reflects
 101 only its prior belief about its own accuracy $\hat{a}$. We then compare $\hat{a}$ to its actual accuracy. To
 102 then assess threshold alignment, we construct escalation curves across accuracy levels and
 103 identify the accuracy p^* at which the agent’s escalation rate crosses 50%. A perfectly aligned
 104 agent at cost ratio R would have $p^* = \tau^* = 1 - 1/R$.

105 3 Experimental design

106 3.1 Datasets

107 We evaluate LLMs on four tasks, each drawn from a large-scale dataset of recorded human
 108 decisions. We also include a fifth dataset (*MoralMachine*) as a robustness check; because
 109 ethical dilemmas are not a typical agent task, we report those results in Appendix A. Each
 110 setting centers around a binary decision for which we have real data on human preference.

111 **Demand forecasting (*HotelBookings*).** We predict whether a hotel booking will be kept or
 112 canceled, using features such as lead time, number of special requests, and guest history
 113 (Antonio et al., 2019). The dataset contains 119,390 bookings.

114 **Loan approval (*LendingClub*).** We predict whether a loan application will be approved,
 115 using applicant features such as FICO score, debt-to-income ratio, and loan amount (George,
 116 2018). We sample 20,000 applications (10,000 accepted, 10,000 rejected) from the full dataset
 117 of approximately 2.26 million loans.

118 **Content moderation (*Wikipedia Toxicity*).** We predict whether a Wikipedia discussion
 119 comment will be labeled toxic by crowd-workers (Wulczyn et al., 2017). The dataset contains
 120 159,686 unique comments.

121 **Content recommendation (*MovieLens*).** We predict which of two movies a user would rate
 122 more highly, given five prior ratings and average community ratings (Harper & Konstan,
 123 2015). We construct pairwise comparisons from a sample of 1,000,000 ratings by 6,307 users
 124 from the full *MovieLens* 25M dataset.

125 ***MoralMachine* (Appendix).** We predict which group an autonomous vehicle should save
 126 in an ethical dilemma, based on the *Moral Machine* dataset, from which we sample 500,000

scenarios from the full 1M-response US survey. The dataset also includes demographic information about the human *driver* (Awad et al., 2018).

Example scenarios for each dataset are provided in Appendix G.

3.2 Models

The main text reports eight models in four families, each with a smaller and larger variant:

- **Qwen3.5-4B / Qwen3.5-9B:** Dense models from the Qwen3.5 family.
- **Gemma 3 4B / Gemma 3 12B:** Dense models from Google.
- **Claude Sonnet 4.6 / Claude Opus 4.7:** Closed-source models from Anthropic.
- **GPT-5-nano / GPT-5-mini:** Closed-source reasoning models from OpenAI.

The Appendix extends the analysis to the MoE flagship Qwen3.5-397B-A17B and two more families: Llama 4 Maverick / Llama 3.3 70B (Meta) and Mixtral 8x7B / Mistral Small 24B (Mistral).

GPT-5-mini excludes *MovieLens* and *MoralMachine* due to content filter restrictions. GPT-5-nano excludes *MoralMachine* for the same reason. Gemma 3 and Claude models are evaluated on the four main datasets but not *MoralMachine*. The remaining models are evaluated on all five datasets.

3.3 Prompting protocol

Each scenario is accompanied by a *signal*: a natural-language summary of a decision tree’s output for the scenario’s feature-defined condition, including the feature split and its predictive accuracy (e.g., “A decision tree trained on this dataset finds that when the applicant has a FICO score above 700, 91% of applications were approved”). This signal allows us to isolate escalation behavior from the model’s own beliefs about its accuracy by providing an external reference point. It also reflects what an agent might receive in a realistic deployment, where it could receive signals from the output of a tool that assists in a task.

Each sample proceeds in two turns within a single conversation. In the first turn, the agent receives the scenario, the signal (when applicable), and a prediction prompt asking it to explain its reasoning in one sentence and conclude with PREDICTION: 1 or PREDICTION: 0. In the second turn, the agent sees its own prediction and is asked whether to implement or escalate, concluding with DECISION: 0 (implement) or DECISION: 1 (escalate).

We evaluate several variants of this protocol. **Baseline** and **No signal** are used in Sections 4 and 5 to characterize each model’s intrinsic escalation behavior. **Cost framing**, **Thinking**, and their combination are interventions evaluated in Section 6.

- **Baseline:** Signal provided, no cost framing, no explicit thinking.
- **No signal:** The signal is omitted, revealing the model’s default confidence.
- **Cost framing** ($R = 4$): The escalation prompt is prepended with “Implementing a wrong answer costs $4\times$ more than escalating.”
- **Thinking:** For Qwen3.5-9B, we enable the model’s extended thinking mode. For GPT-5-mini, we increase reasoning_effort from minimal to medium.
- **Thinking + cost framing:** Both interventions are active.

For each condition, we sample 250 scenarios (100 for thinking runs) and report means with standard errors.

3.4 Metrics

We use three accuracy measures throughout, and refer to each by its specific name to avoid ambiguity.




Figure 2: The multi-turn prompting protocol, illustrated with a *LendingClub* example. In Turn 1, the agent receives the scenario and signal, then produces a prediction. In Turn 2, the agent sees its own prediction and decides whether to implement or escalate. The signal provides the feature split and predictive accuracy from a decision tree trained on the dataset.

Predictive accuracy (p) is the rate at which the agent's prediction $\hat{y}$ matches the ground truth y within a condition. It measures how well the agent's prediction tracks the underlying labels and depends only on the prediction in Turn 1.

Escalation accuracy is the rate at which the agent's implement/escalate choice in Turn 2 agrees with the cost-optimal action for a given cost ratio R . An escalation is optimal when the prediction's expected error cost exceeds the labor cost, i.e., when $p < \tau^* = 1 - 1/R$. Escalation accuracy is therefore a property of the joint behavior across both turns, evaluated against the Bayes-optimal policy.

Self-estimated predictive accuracy ($\hat{a}$) is the agent's implicit belief about its own predictive accuracy, recovered by mapping its no-signal escalation rate onto the signal-based escalation curve (Section 5).

4 Escalation profiles

Models form internal beliefs about their own predictive accuracy, and these beliefs shape their escalation decisions. To isolate escalation behavior from these beliefs, we provide an explicit signal that specifies the predictive accuracy for each scenario. Figure 3 presents the resulting escalation rate versus predictive accuracy for each model across all four datasets. Several patterns emerge.

First, every model escalates less as its predictive accuracy increases, so all models are at least somewhat sensitive to the signal. Second, however, the models differ dramatically in their overall escalation levels: the same accuracy can produce vastly different escalation rates depending on the model. Third, several models' curves are non-monotonic, suggesting that models incorporate factors beyond accuracy when deciding whether to escalate, or that escalation behavior is inherently noisy.

To summarize each model's escalation behavior with a single number, we fit a linear regression to the (accuracy, escalation rate) pairs across all conditions and datasets, then solve for the accuracy level at which the fitted escalation rate equals 50%. We call this the model's *implicit threshold* p^* .

Figure 6 (Appendix E) summarizes these values and reveals striking variation. Across our eight main-text models, p^* ranges from 53% (GPT-5-mini) to 98% (Gemma 3 12B), spanning nearly the full plausible range. Scaling within a family shifts behavior unpredictably, and not only in threshold magnitude. Qwen3.5-4B ($p^* = 61\%$) and Qwen3.5-9B ($p^* = 56\%$) sit close in threshold but differ sharply in the *shape* of their escalation curves: 4B approximates a step function, escalating almost always below its threshold and almost never above, while 9B responds gradually and non-monotonically. Gemma 3 shows a large threshold gap: 4B sits at $p^* = 76\%$ while 12B jumps to $p^* = 98\%$. GPT-5-nano ($p^* = 91\%$) and GPT-5-mini ($p^* = 53\%$) differ by 38 points. The Claude family is more tightly clustered, with Sonnet at




Figure 3: Each model exhibits a distinct escalation profile. Escalation rate versus predictive accuracy for all eight models across four datasets (with signals, no thinking). Same-family models are vertically aligned. Models vary in both overall escalation level and sensitivity to the signal. Error bars show ± 1 standard error.

208 $p^* = 66\%$ and Opus at $p^* = 59\%$, both leaning toward implementation. Appendix F reports
 209 the same analysis for five additional models, including the Qwen3.5-397B MoE flagship
 210 ($p^* = 81\%$, a 25-point gap above the dense 9B) and the Llama and Mistral families, where
 211 within-family gaps of 30 to 40 percentage points also recur.

212 For a practitioner choosing a model for automated decision-making, these differences are
 213 consequential. A model with $p^* = 56\%$ will implement aggressively, accepting frequent
 214 errors in exchange for fewer escalations. A model with $p^* = 91\%$ will escalate most
 215 decisions, providing safety at the expense of automation. Neither behavior is inherently
 216 correct; the optimal threshold depends on the cost ratio R . These thresholds are latent,
 217 model-specific properties that cannot be determined without empirical characterization.

218 5 Self-estimated predictive accuracy

219 Having established signal-based escalation curves, we now recover what a model believes
 220 about its own predictive accuracy. Without a signal, the model escalates at a rate that reflects
 221 only its prior beliefs. We map this no-signal escalation rate onto the signal-based curve
 222 to find the predictive accuracy that would produce the same escalation rate, yielding the
 223 model’s *self-estimated predictive accuracy* $\hat{a}$.

224 A substantial literature documents miscalibration in LLMs, predominantly framing it as
 225 systematic overconfidence on factual recall and question answering (Kadavath et al., 2022;
 226 Xiong et al., 2024; Tian et al., 2023; Lin et al., 2022; Mielke et al., 2022; Guo et al., 2017).
 227 Our results extend this picture along two axes. First, miscalibration in escalation behavior
 228 is not unidirectional: models span the full range from systematically overconfident to
 229 systematically underconfident, with several models nearly calibrated. Second, the direction
 230 is model-specific and unpredictable from architecture or scale.

231 We find that calibration varies markedly across our main-text models (Figure 4). At the
 232 overconfident end, Qwen3.5-9B exhibits a 19-point gap between its self-estimated and actual
 233 predictive accuracy (94% vs. 75%); GPT-5-nano shows moderate overconfidence (85% vs.
 234 76%). At the underconfident end, both Gemma 3 sizes underestimate their accuracy by
 235 7–8 points (self-estimates of 68–69% against actual 75–76%). Three models, Claude Son-
 236 net 4.6, Claude Opus 4.7, and GPT-5-mini, sit close to calibrated, with self-estimates within
 237 3 percentage points of actual. Within the same family, scaling up shifts calibration unpre-
 238 dictably: dense Qwen3.5 moves from mild overconfidence at 4B to severe overconfidence




Figure 4: Actual versus self-estimated predictive accuracy for each of the eight main-text models across four datasets. Points above the diagonal indicate overconfidence (self-estimated > actual). The direction and magnitude of miscalibration vary markedly by model: Qwen3.5-9B is consistently overconfident, Gemma 3 is underconfident, GPT-5-nano is moderately overconfident, and Claude and GPT-5-mini are close to calibrated. Appendix F shows the same plot for the additional five models.

at 9B, GPT-5 moves from overconfident at nano to calibrated at mini, and Gemma 3 stays underconfident at both sizes. Appendix F extends this analysis to five additional models, including Qwen3.5-37B and the Llama and Mistral families.

These aggregate numbers, however, understate the problem (Figure 6, Appendix E). Self-estimates span 68% (Gemma 3 4B) to 94% (Qwen3.5-9B), while actual average predictive accuracy stays narrowly within 75% to 80%. Qwen3.5-9B exemplifies severe overconfidence: it is overconfident on nearly every condition, with gaps ranging from -2 to $+41$ percentage points. Aggregate calibration can also mask per-condition variation: Claude Sonnet appears nearly calibrated overall, yet its condition-level gaps still span both directions. Appendix F discusses the additional models, including cases where offsetting biases hide miscalibration in the aggregate.

Even knowing a model’s self-estimated predictive accuracy does not predict its escalation threshold. Self-estimated predictive accuracy spans 68% to 94%, and p^* spans 53% to 98%, but the two are largely independent. Gemma 3 12B is underconfident yet has the highest threshold (escalates aggressively), while GPT-5-mini is nearly calibrated yet has the lowest threshold (implements aggressively). Practitioners must characterize both dimensions before deployment.

6 Calibrating the threshold

We now test whether the escalation threshold can be aligned with a target cost ratio through prompting interventions and fine-tuning.

6.1 Prompting interventions

We evaluate Qwen3.5-9B under four conditions: baseline (signal, no thinking, no cost framing), cost framing alone, thinking alone, and thinking with cost framing. For each sample, we score whether the model makes the cost-optimal escalation decision at $R = 4$, which sets the optimal threshold at $\tau^* = 75\%$.




Figure 5: Cost framing intervention. The cost statement (red box) is prepended to the Turn 2 user message, exposing the cost ratio $R = 4$ to the agent. All other prompt elements are unchanged from the baseline protocol (Figure 2).

Table 1: Sample-level escalation accuracy under different interventions, evaluated against the optimal policy for cost ratio $R = 4$ ($\tau^* = 75\%$, $c_w = 4c_\ell$). For each sample, the correct action is to escalate if the condition’s predictive accuracy is below 75% and implement if above. 95% intervals are Wilson score intervals on the binomial proportion of correct decisions (reproduced by scripts/compute_intervention_cis.py). *MoralMachine* is excluded from all rows. GPT-5-mini further excludes *MovieLens* due to content filter restrictions.

Intervention	Accuracy (95% CI)	Δ
Qwen 9B baseline	59.0% [58.3, 59.7]	—
+ cost framing	59.8% [59.1, 60.5]	+0.8
+ thinking	58.1% [56.7, 59.4]	−0.9
+ thinking + cost framing	80.5% [79.4, 81.5]	+21.5
GPT-5-mini baseline	66.3% [65.3, 67.4]	—
+ cost framing	77.7% [76.7, 78.6]	+11.4
+ reasoning	71.2% [70.1, 72.3]	+4.9
+ reasoning + cost framing	91.4% [90.4, 92.3]	+25.1

Table 1 reveals a clear interaction effect. Cost framing alone barely improves escalation accuracy for Qwen3.5-9B (59.8% vs. 59.0% baseline). Thinking alone is slightly worse than baseline (58.1%), because the model’s predictive accuracy rises but its escalation rate drops, leading to over-implementation.

The combination of thinking and cost framing achieves 80.5% escalation accuracy, a substantial improvement over all other Qwen conditions. Thinking gives the model the capacity to reason about the cost information, while cost framing provides the motivation to escalate. Neither intervention is sufficient alone; they are complementary.

GPT-5-mini uses built-in reasoning controlled by OpenAI’s reasoning_effort parameter. The lowest available setting is minimal (OpenAI does not support disabling reasoning entirely), so some residual reasoning may contribute to the baseline. At the minimal level, it achieves 66.3% escalation accuracy, comparable to the Qwen baseline. Unlike Qwen, cost framing alone improves GPT-5-mini to 77.7% even at minimal reasoning, suggesting that the model retains some capacity to process cost information without extended reasoning. At medium reasoning, escalation accuracy reaches 71.2%, and the combination of medium reasoning and cost framing achieves 91.4%.

6.2 Fine-tuning for cost-sensitive escalation

Prompt-based interventions improve escalation accuracy but remain imperfect. We investigate whether fine-tuning can train a model to internalize cost-sensitive escalation behavior. We train three SFT models with chain-of-thought responses that explicitly ex-

tract the predictive accuracy from the signal and compute the expected cost: Qwen2.5-7B-Instruct, Qwen3.5-4B, and Qwen3.5-9B. The target response follows a template: “The signal suggests a predictive accuracy of $X\%$, so the error rate is $Y\%$. The expected cost of implementing is $R \times Y = Z$, which exceeds/is below 1.” We train on three of the four main datasets (*HotelBookings*, *LendingClub*, *WikipediaToxicity*), four cost framings, and six cost ratios ($R \in \{2, 4, 8, 10, 20, 50\}$), holding out *MovieLens* entirely.

All three SFT models substantially outperform the prompt-only interventions in Table 1. Qwen2.5-7B-Instruct achieves near-perfect escalation accuracy: 100% on most cells with a single error at a boundary case where $R(1 - p) = 1.00$ exactly. Qwen3.5-4B reaches 95.7% overall (held-out *MovieLens* 96.7%), and Qwen3.5-9B reaches 87.3% overall (held-out 89.2%). The held-out generalization across all three models indicates that SFT teaches a general procedure rather than memorizing dataset-specific decisions. The within-Qwen3.5 ranking is non-monotonic, with 4B exceeding 9B, suggesting the same recipe may benefit from per-model tuning. Full per-dataset and per-framing results appear in Table 4, Appendix D.

To confirm that the SFT model reads and uses the signal, we run an ablation that removes it entirely. Without the signal, escalation accuracy drops from 100% to approximately 84% as the model hallucinates predictive-accuracy estimates instead of reading them from the signal (Appendix B). The chain-of-thought supervision is critical: SFT directly supervises a reasoning procedure that extracts the predictive accuracy from the signal and computes expected cost, enabling generalization across datasets, cost framings, and held-out domains.

7 Related work

Our work connects to LLM calibration, where Kadavath et al. (2022) show that language models can assess their own uncertainty and Xiong et al. (2024) evaluate confidence elicitation methods. Guo et al. (2017) demonstrate that modern neural networks are poorly calibrated, and Tian et al. (2023) find that verbalized confidence from LLMs can outperform token-level probabilities in some settings. We build on this line of work by asking how self-assessed uncertainty affects the downstream decision to act or defer, rather than evaluating confidence scores alone.

In LLM agents, deferral can be analyzed as a threshold decision, but the deployed policy is usually not exposed as a separate classifier threshold that a practitioner can inspect or tune directly. It is elicited through language and can vary across prompts and models. We therefore draw on selective prediction (Geifman & El-Yaniv, 2017), and the related literature on learning to defer, where a classifier can route instances to a human expert (Madras et al., 2018; Mozannar & Sontag, 2020). Our cost-based formulation connects to cost-sensitive classification (Elkan, 2001), which optimizes decision thresholds under asymmetric misclassification costs. We do not propose a new abstention theory. Instead, we adapt these ideas to deployed agent behavior, where the act-or-defer policy is part of the model’s prompted behavior rather than a separately parameterized decision rule.

Chain-of-thought prompting (Wei et al., 2022) enables models to reason step-by-step, improving performance on arithmetic and multi-step tasks. Our finding that extended thinking is necessary for cost framing to take effect aligns with evidence that explicit reasoning unlocks capabilities latent in the base model. On the alignment side, RLHF (Ouyang et al., 2022) and DPO (Rafailov et al., 2023) train models to match human preferences; we use SFT with chain-of-thought supervision to align escalation behavior with the optimal policy.

Our multi-turn prompting protocol is related to multi-agent debate (Du et al., 2023; Liang et al., 2023), but the second turn evaluates the agent’s own prediction rather than introducing an adversarial perspective.

8 Conclusion

We have shown that LLM-based agents exhibit latent, model-specific escalation profiles that vary dramatically across models and cannot be predicted *a priori* from architecture or scale alone. Even within the same model family, scaling up produces unpredictable

shifts: Gemma 3 4B and 12B differ by 22 percentage points in their implicit thresholds, and GPT-5-nano and GPT-5-mini differ by 38 points. Scaling dense Qwen3.5 from 4B to 9B barely moves the threshold but sharply worsens calibration and turns a step-like escalation curve into a noisy, gradual one. Models also carry miscalibrated beliefs about their own predictive accuracy, and the direction of miscalibration can reverse between a smaller and larger variant of the same family. Even calibration itself is model-specific: Claude Sonnet and Opus closely match their actual predictive accuracy, GPT-5-mini lands close to calibrated, while Gemma 3 models systematically underestimate it. Appendix F confirms these patterns extend to five additional models, including the Qwen3.5-397B flagship and the Llama and Mistral families. These latent decisions represent below-the-surface behavior that could disrupt a decision-making process without sufficient exploration.

While most LLM evaluation focuses on *capability* (accuracy, throughput, cost-savings), our results characterize a related deployment behavior that accuracy alone does not measure: how an agent *decides* to act or defer under uncertainty. This behavior is latent, model-specific, and not predicted by architecture or scale, but as we show it can be characterized empirically and corrected through targeted training. Capability-focused benchmarks miss this failure mode: a high-accuracy model can still systematically act when it should defer or defer when it should act, undermining the automation it was deployed for.

These findings have immediate and extensive practical implications. Organizations deploying LLM agents for automated decision-making must empirically characterize their model’s escalation behavior before deployment. Our methodology, based on varying predictive accuracy through calibrated signals, provides a tractable approach for doing so.

We also demonstrate that escalation behavior can be corrected through appropriate interventions. The combination of extended thinking and cost framing achieves near-optimal escalation decisions, and supervised fine-tuning pushes escalation accuracy even higher while generalizing across domains, framings, and costs.

Several limitations suggest directions for future work. Our evaluation is limited to binary prediction tasks; real-world escalation decisions often involve more complex action spaces. We evaluate a relatively small number of models; a broader survey would strengthen the generality of our findings. Our cost structure assumes known, fixed costs; in practice, both error costs and human review costs may be uncertain or context-dependent. We also model escalation as deferral to the labeling authority. In deployed human-in-the-loop systems, human review may be imperfect because automation can shift work, reduce context, or create new coordination burdens (Moradi et al., 2024). Our results therefore characterize the agent-side act-or-defer decision, not the full sociotechnical cost of human fallback.

Ethics Statement. Our experiments use public, de-identified datasets of recorded human decisions and involve no human subjects and no new data collection. Some datasets cover consequential domains, including loan approval and content moderation, which we use only to study when a model defers, not as endorsements of automating those decisions.

LLM Disclosure. Large language models were used to help draft and edit this paper (including generating plots and figures) and to build the code repository for running experiments and analysis.

Data and Code Availability. All code, data, and experimental results are publicly available at <https://github.com/anonymous-code-submitter/colm-2026-escalation>.

References

- Nuno Antonio, Ana de Almeida, and Luis Nunes. Hotel booking demand datasets. *Data in Brief*, 22:41–49, 2019. doi: 10.1016/j.dib.2018.11.126.
- Edmond Awad, Sohan Dsouza, Richard Kim, Jonathan Schulz, Joseph Henrich, Azim Shariff, Jean-François Bonnefon, and Iyad Rahwan. The moral machine experiment. *Nature*, 563: 59–64, 2018. doi: 10.1038/s41586-018-0637-6.
- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot arena: An open platform for evaluating LLMs by human preference. *arXiv preprint arXiv:2403.04132*, 2024.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.
- Charles Elkan. The foundations of cost-sensitive learning. In *International Joint Conference on Artificial Intelligence*, pp. 973–978, 2001.
- Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems*, 2017.
- Nathan George. Lending club loan data. <https://www.kaggle.com/datasets/wordsforthewise/lending-club>, 2018.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning*, pp. 1321–1330, 2017.
- F. Maxwell Harper and Joseph A. Konstan. The MovieLens datasets: History and context. *ACM Transactions on Interactive Intelligent Systems*, 5(4):19:1–19:19, 2015. doi: 10.1145/2827872.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2024.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *arXiv preprint arXiv:2305.19118*, 2023.
- Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*, 2022.
- David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems*, 2018.
- Sabrina J. Mielke, Arthur Szlam, Emily Dinan, and Y-Lan Boureau. Reducing conversational agents’ overconfidence through linguistic calibration. *Transactions of the Association for Computational Linguistics*, 10:857–872, 2022.
- Pegah Moradi, Karen Levy, and Cristobal Cheyre. Pseudo-automation: How labor-offsetting technologies reconfigure roles and relationships in frontline retail work, 2024.
- Hussein Mozannar and David Sontag. Consistent estimators for learning to defer to an expert. In *International Conference on Machine Learning*, pp. 7076–7087, 2020.

- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2023.
- Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, and Chelsea Finn. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. *arXiv preprint arXiv:2305.14975*, 2023.
- Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin Manku, Vinty Dong, Edward Li, Shashank Gupta, Ashish Sabharwal, and Niranjana Balasubramanian. AppWorld: A controllable world of apps and people for benchmarking interactive coding agents. *arXiv preprint arXiv:2407.18901*, 2024.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, et al. A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6), 2024. doi: 10.1007/s11704-024-40231-1.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. AutoGen: Enabling next-gen LLM applications via multi-agent conversation. In *COLM*, 2024.
- Ellery Wulczyn, Nithum Thain, and Lucas Dixon. Ex machina: Personal attacks seen at scale. In *Proceedings of the 26th International Conference on World Wide Web*, pp. 1391–1399, 2017. doi: 10.1145/3038912.3052591.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. *arXiv preprint arXiv:2306.13063*, 2024.

A *MoralMachine* results

MoralMachine presents ethical dilemmas rather than typical agent tasks, making it a robustness check for whether escalation patterns generalize beyond standard decision-making. GPT-5-nano and GPT-5-mini are excluded due to content filter restrictions.

Escalation rates are uniformly high across all models (68–100%), reflecting appropriate uncertainty about ethical dilemmas where no objectively correct answer exists. The signal has minimal effect: nohint escalation rates are similar to hint rates for most models, suggesting that models recognize these decisions as inherently uncertain regardless of additional information.

B SFT signal ablation

To verify that the SFT model relies on the feature signal rather than exploiting some other signal, we evaluate it on prompts that omit the signal entirely. Table 3 reports the results.

Without the signal, escalation accuracy drops from 100% to approximately 84%. The model still follows the trained reasoning template, but it hallucinates a predictive-accuracy estimate (typically 75–85%) rather than reading it from the prompt. The hallucinated estimates are systematically optimistic, which produces correct escalation decisions at extreme cost ratios

Table 2: *MoralMachine* escalation behavior. All models escalate at high rates, recognizing the inherent uncertainty of ethical dilemmas.

Model	With signal		Without signal	
	Pred. acc	Esc. rate	Pred. acc	Esc. rate
Qwen3.5-9B	61.2%	68.3%	52.1%	70.6%
Qwen3.5-397B	70.4%	93.4%	69.3%	91.3%
Llama 4 Maverick	65.8%	84.0%	65.2%	94.6%
Llama 3.3 70B	67.5%	97.7%	71.4%	99.9%
Mixtral 8x7B	65.1%	96.5%	39.5%	93.5%
Mistral Small 24B	70.2%	98.2%	68.4%	100.0%

Table 3: SFT model evaluated with and without the signal, on the original cost framing.

Dataset	With signal	Without signal
<i>HotelBookings</i>	100%	84.7%
<i>LendingClub</i>	100%	84.0%

(where the decision is insensitive to the exact predictive accuracy) but incorrect decisions at $R = 4$ (where escalation accuracy drops to 48–68%). The model’s arithmetic remains correct throughout; only the input is wrong. This confirms that the SFT model genuinely reads and uses the feature signal, and that the signal is the critical ingredient for perfect performance.

C Proofs

Proof of Theorem 1. For a single instance with true accuracy p , the expected cost under threshold τ is

$$C(p, \tau) = \begin{cases} c_\ell & \text{if } p < \tau \text{ (escalate)} \\ (1 - p) c_w & \text{if } p \geq \tau \text{ (implement).} \end{cases}$$

The agent should escalate when $c_\ell < (1 - p) c_w$, i.e., when $p < 1 - c_\ell / c_w = \tau^*$. This threshold uniquely minimizes expected cost because it is the only value at which the agent is indifferent between escalating and implementing. $\square$

D SFT results

E Escalation threshold and self-estimated predictive accuracy

Figure 6 summarizes the implicit threshold p^* and self-estimated predictive accuracy $\hat{a}$ for all thirteen models, as described in Sections 4 and 5.

F Additional models

The main text reports eight models in four families (Qwen3.5, Gemma 3, Claude, GPT-5). We additionally evaluated five models using the same protocol: the Qwen3.5-397B-A17B MoE flagship, plus Llama 4 Maverick and Llama 3.3 70B (Meta), and Mixtral 8x7B and Mistral Small 24B (Mistral). The escalation-threshold and self-calibration patterns established in the main text extend to these models, with comparable within-family variation. Figure 6 reports the implicit threshold p^* and self-estimated predictive accuracy $\hat{a}$ for all thirteen models side by side.

Qwen3.5-397B has $p^* = 81\%$, 25 percentage points above the dense Qwen3.5-9B ($p^* = 56\%$), the largest within-Qwen threshold gap and a useful contrast to the near-flat 4B-to-9B threshold shift in the main text. Among the Llama and Mistral models, p^* exceeds 90%

Table 4: SFT escalation accuracy for three fine-tuned models. Each is trained on three datasets and evaluated on all four main datasets, including the held-out *MovieLens*. Accuracy is measured against the Bayes-optimal policy across six cost ratios ($R \in \{2, 4, 8, 10, 20, 50\}$), with 60 samples per (dataset, framing). Cells marked — were not evaluated (Qwen2.5-7B-Instruct predates the addition of the Study3 framing). *MoralMachine* results are reported in Appendix A.

Model	Dataset	Original	Dollar	Wording	Study3
Qwen2.5-7B-Instruct	<i>HotelBookings</i> (train)	100%	98.3%	100%	—
	<i>LendingClub</i> (train)	100%	100%	100%	—
	<i>WikipediaToxicity</i> (train)	100%	100%	100%	—
	<i>MovieLens</i> (held-out)	100%	100%	100%	—
Qwen3.5-4B	<i>HotelBookings</i> (train)	98.3%	88.3%	98.3%	98.3%
	<i>LendingClub</i> (train)	95.0%	95.0%	95.0%	90.0%
	<i>WikipediaToxicity</i> (train)	96.7%	98.3%	95.0%	96.7%
	<i>MovieLens</i> (held-out)	98.3%	91.7%	98.3%	98.3%
Qwen3.5-9B	<i>HotelBookings</i> (train)	91.7%	98.3%	90.0%	88.3%
	<i>LendingClub</i> (train)	93.3%	90.0%	86.7%	83.3%
	<i>WikipediaToxicity</i> (train)	76.7%	93.3%	73.3%	75.0%
	<i>MovieLens</i> (held-out)	93.3%	85.0%	85.0%	93.3%




Figure 6: Implicit escalation threshold p^* (left) and self-estimated predictive accuracy $\hat{a}$ (right) for each of the thirteen models. The threshold p^* spans 56% (Qwen3.5-9B) to nearly 100% (high-threshold models), while self-estimated predictive accuracy ranges from 68% to 97%. The dotted lines show the optimal threshold $\tau^* = 75\%$ at cost ratio $R = 4$ (left) and average actual accuracy (right).

for all four, with Llama 3.3 70B and Mixtral 8x7B above 100%, and Mixtral 8x7B and Mistral Small 24B differ by more than 30 percentage points within the Mistral family. On the calibration axis, Mixtral 8x7B and Mistral Small 24B are strongly overconfident, Llama 3.3 70B is underconfident on aggregate but with offsetting per-condition biases, and Llama 4 Maverick falls in between. These results corroborate the main-text finding that escalation behavior is model-specific and unpredictable from architecture or scale.

G Example scenarios

HotelBookings. “Person 1 has booked a hotel stay arriving on July 12, 2017 (week 28), with 2 weekend night(s) and 3 weekday night(s). The party consists of 2 adult(s). Person 1 is not a repeated guest and has 0 previous cancellation(s). They have requested 1 car parking space(s) and made 2 special request(s).”

506 **LendingClub.** “The applicant is requesting \$12,000 for debt consolidation. They have been
507 employed for 5 years, a debt-to-income ratio of 14.3%, and a credit score of 712.”

508 **Wikipedia Toxicity.** “This comment needs to be checked: ‘Please stop removing content
509 from Wikipedia. It is considered vandalism.’”

510 **MovieLens.** “Person 1 has reviewed: Toy Story (4/5), Braveheart (5/5), Apollo 13 (4/5),
511 Pulp Fiction (5/5), Forrest Gump (3/5). Consider these two movies: Movie A: The Shaw-
512 shank Redemption (Drama), average rating 4.43/5 (311 ratings). Movie B: Twelve Monkeys
513 (Mystery—Sci-Fi—Thriller), average rating 3.64/5 (229 ratings).”

514 **MoralMachine.** “An autonomous vehicle is about to get in an accident. If the car doesn’t
515 swerve, 2 elderly pedestrians will die. If the car swerves, 3 passengers will die. The
516 pedestrians are legally crossing the street. The person behind the wheel is a 34-year-old
517 male with a Bachelor Degree.”

518 H Implementation details

519 This appendix lists the hyperparameters and inference settings used in our experiments. All
520 code is also available in the project repository.

521 **Inference (Sections 4–6).** All models are queried through an OpenAI-compatible chat-
522 completions interface. Decoding is greedy (temperature = 0) and no sampling parameters
523 (top_p, top_k, etc.) are altered from provider defaults. Per-turn generation budgets:

- 524 • **Prediction turn** (Turn 1): max_tokens = 256 for non-reasoning models. For OpenAI
525 reasoning models (GPT-5-nano, GPT-5-mini), max_completion_tokens = 2048 to
526 accommodate reasoning tokens.
- 527 • **Escalation turn** (Turn 2): same per-model settings as Turn 1, except in extended-
528 thinking runs where max_tokens = 16,384 to leave room for the model’s reasoning
529 trace.
- 530 • **Together client:** timeout = 150 s per request, with up to 5 retries on transient
531 failures (exponential backoff).
- 532 • **Qwen3.5 thinking:** enabled by leaving the default chat-template behavior on;
533 disabled by passing extra_body = {"chat_template_kwargs": {"enable_thinking":
534 False}}.
- 535 • **GPT-5 reasoning:** the “minimal” setting in Table 1 corresponds to reasoning_effort
536 = “minimal”; “medium” uses the default reasoning_effort = “medium”.

537 For each condition we sample 250 scenarios (100 for extended-thinking runs) using
538 random_state = 43, and we report means with standard errors.

539 **SFT training (Section 6).** All three SFT models use the same recipe
540 (scripts/sft_escalation.py). LoRA configuration: $r = 64$, $\alpha = 128$, target mod-
541 ules = all-linear, dropout = 0.05, bias = none, task type = CAUSAL_LM. SFTTrainer
542 settings: 1 epoch, per_device_train_batch_size = 1, gradient_accumulation_steps =
543 8 (effective batch size 8), learning rate 5×10^{-5} with default linear warmup, bf16 =
544 True, gradient_checkpointing = True, max_length = 768. Training data is constructed
545 from 10 samples per condition across 3 training datasets (*HotelBookings*, *LendingClub*,
546 *WikipediaToxicity*), expanded across 4 cost framings and 6 cost ratios ($R \in \{2, 4, 8, 10, 20, 50\}$),
547 yielding $\sim 7,200$ training examples. *MovieLens* is held out entirely.

548 **SFT evaluation (Section 6, Appendix D).** The fine-tuned model is loaded with its LoRA
549 adapter via peft.PeftModel.from_pretrained, in bfloat16. Generation is greedy with
550 do_sample = False, max_new_tokens = 128 for Turn 1 and 256 for Turn 2. The chat tem-
551 plate is applied with enable_thinking = False to match the training setup. We parse the

model’s decision from the *first* DECISION: 0/1 token in the response; the model occasionally continues generating fictitious follow-up turns whose later decisions are unreliable.

Hardware. Inference for non-fine-tuned models runs through provider APIs (no local hardware). SFT training and evaluation run on AMD MI210/MI300 GPUs via the AMD HPC Fund allocation. Each SFT training run uses one GPU and completes in 1.5–3.5 h depending on model size; each per-dataset SFT-evaluation run completes in 1–3 h.

I Qualitative model responses

We show four verbatim model responses on the same *HotelBookings* condition (NO_PREV_CANCEL, base rate 66%, $R = 4$ where applicable). At this condition the optimal action is to *escalate* (signal accuracy below $\tau^* = 75\%$), so a correct decision is DECISION: 1. The four responses reach *different* conclusions despite seeing the same signal, illustrating model-specific calibration and the effect of interventions.

Qwen3.5-9B baseline (no intervention). Implements with a tautological justification, never engaging with the signal’s accuracy:

“Since my previous prediction was that the booking would be kept (Prediction: 1), implementing the current reservation without escalation is the appropriate course of action to align with the predicted outcome. DECISION: 0”

The model treats its own prediction as license to act, with no cost reasoning at all. **Incorrect.**

Claude Opus 4.7 baseline. Acknowledges the uncertainty but still acts:

“The prediction is based on a relatively weak signal (66% retention rate), meaning there’s still a ~34% chance of being wrong, but without additional risk factors, implementing aligns with the most likely outcome. Given the moderate confidence and no strong red flags, implementing is reasonable. DECISION: 0”

The model reads the signal correctly but resolves the uncertainty in favor of acting. **Incorrect.**

Claude Sonnet 4.6 baseline. Reads the *same* 66%/34% numbers and reaches the opposite conclusion:

“The prediction is based on a moderate 66% confidence from the decision tree, which means there is a substantial 34% chance of being wrong. Given this relatively low confidence level, it would be prudent to escalate to Person 1 to confirm their booking intentions rather than risk an incorrect implementation. DECISION: 1”

A within-family split on identical signal: Opus implements, Sonnet escalates. **Correct.**

Qwen3.5-9B with thinking and cost framing. The same Qwen that was tautological at baseline now does the cost-benefit reasoning explicitly:

“Given the 66% historical retention rate implies a 34% cancellation risk, and implementing an error incurs a four-fold higher penalty than escalation, the risk of a wrong implementation outweighs the direct cost of human intervention. DECISION: 1”

The intervention flips the decision via explicit cost arithmetic, despite the same signal. **Correct.**

The four responses illustrate the paper’s three main findings in one comparable case. First, baseline calibration is model-specific: the same 66%/34% numbers lead Opus and Sonnet to opposite decisions within the Claude family, and lead Qwen3.5-9B to ignore the signal altogether. Second, the choice to act under uncertainty is not predicted by self-stated

595 reasoning: Opus and Qwen3.5-9B both implement, but for very different reasons. Third,
596 interventions repair Qwen3.5-9B's overconfidence on the same prompt: thinking plus a
597 stated cost ratio produces the correct expected-cost calculation and flips the decision.