

Act or Escalate? Evaluating Escalation Behavior in Automation with Language Models

Anonymous authors

Paper under double-blind review

Abstract

Effective automation hinges on deciding when to act and when to escalate. We model this as a decision under uncertainty: an LLM forms a prediction, estimates its probability of being correct, and compares the expected costs of acting and escalating. Using this framework across five domains of recorded human decisions—demand forecasting, content recommendation, content moderation, loan approval, and autonomous driving—and across multiple model families, we find marked differences in the implicit thresholds models use to trade off these costs. These thresholds vary substantially and are not predicted by architecture or scale, while self-estimates are miscalibrated in model-specific ways. We then test interventions that target this decision process by varying cost ratios, providing explicit accuracy hints, and training models to follow the desired escalation rule. Prompting helps mainly for reasoning models. SFT on chain-of-thought targets yields the most robust policies, which generalize across datasets, cost ratios, prompt framings, and held-out domains. These results suggest that escalation behavior is a model-specific property that should be characterized before deployment, and that robust alignment benefits from training models to reason explicitly about uncertainty and decision costs.

1 Introduction

Agents based on large language models (LLMs) are increasingly deployed to automate consequential decisions, from triaging customer support tickets (Wang et al., 2024) to generating code (Jimenez et al., 2024) and managing enterprise workflows (Wu et al., 2024). Most evaluations focus on speed, accuracy, and cost-savings (Jimenez et al., 2024; Chiang et al., 2024; Trivedi et al., 2024). However, in each setting, an agent faces a choice that receives far less scrutiny: should it implement its own decision, or escalate to a human?

Agent escalation behavior is fundamental to effective automation. An agent that fails to escalate when it is uncertain or incorrect will propagate errors at scale (Lanham et al., 2023), while an agent that always escalates never reduces the human workload it was meant to replace. Two parameters govern effective escalation. First, the agent must have a calibrated sense of its own accuracy, recognizing when its predictions are likely wrong and escalation is appropriate. Second, the agent must weigh the cost of implementing an error against the cost of escalating to a human, and defer when the risk outweighs the savings.

We study escalation behavior across eight models spanning four model families, each represented by a smaller and larger variant: Qwen3.5-9B and Qwen3.5-397B-A17B, GPT-5-nano and GPT-5-mini, Llama 4 Maverick and Llama 3.3 70B, and Mixtral 8x7B and Mistral Small 24B. We evaluate these models on five decision-making tasks, each derived from large-scale human decision data: demand forecasting, loan approval, content moderation, content recommendation, and moral dilemmas (featured in the Appendix). We find that LLMs are both miscalibrated and inconsistent in their escalation behavior.

First, LLMs are miscalibrated in their self-assessment. Some models systematically overestimate their accuracy while others underestimate it, and the direction of miscalibration varies by model and domain. Prior work has documented LLM overconfidence in factual knowledge (Kadavath et al., 2022) and calibration failures in question answering (Xiong




Figure 1: The escalation decision. A human delegates a task to an LLM agent, which produces a prediction and then decides whether to implement (incurring error cost c_w if wrong) or escalate (incurring labor cost c_ℓ). The agent escalates when its estimated probability of being correct $\hat{p}$ falls below a threshold τ .

et al., 2024), but these studies measure confidence in isolation. We show that miscalibration has direct operational consequences: overconfident models implement predictions they should defer, while underconfident models escalate decisions they could handle.

Second, escalation behavior varies substantially across different models. This variance is not associated with model architecture or size. Some models tend to prefer escalating (lower decision threshold), some generally favor implementing (higher decision threshold). This variation persists even within model families. Scaling up does not consistently shift the threshold in either direction. Together, this miscalibration and variance in escalation behavior constitute hidden model dynamics that could disrupt a larger workflow.

These dynamics, however, can be corrected. We test a series of interventions that target the escalation decision. Prompt-based cost framing alone has no effect, but combining it with extended thinking produces improvement for reasoning models. Supervised fine-tuning (SFT) on chain-of-thought targets proves most effective: the resulting model makes near-optimal escalation decisions across all datasets, cost ratios, and prompt framings, and generalizes to held-out domains it was never trained on. Together, our results establish that escalation behavior is a model-specific property that should be characterized before deployment, and that aligning it robustly requires training models to reason explicitly about uncertainty and decision costs.

2 Theory

2.1 The Escalation Decision

Consider a workflow in which a human delegates a binary task to an LLM-based agent. The agent observes a scenario x and produces a prediction $\hat{y} \in \{0, 1\}$. The ground truth is the human’s preference $y \in \{0, 1\}$. The agent then decides whether to *implement* (act on) its prediction or *escalate* (defer to the human). Every action an agent takes in practice contains an implicit escalation decision: when it makes a tool call, generates a response, or commits a change, it is choosing to implement rather than escalate and ask for guidance. Figure 1 illustrates this workflow.

We model the agent’s behavior as follows. After producing its prediction, the agent forms an internal estimate $\hat{p}$ of its probability of being correct $P(\hat{y} = y)$. It then compares $\hat{p}$ to a decision threshold τ (derived below) and escalates if $\hat{p} < \tau$.

2.2 Cost Structure

One of two costs can be incurred depending on the agent’s action. If the agent escalates, the human pays a labor cost $c_\ell > 0$, which represents the time and effort required of the human. If the agent implements and its prediction is wrong, the human pays an error cost $c_w > c_\ell$. Correct implementations incur zero cost.

Theorem 1 (Uniqueness of the optimal threshold). *Let the agent escalate whenever $\hat{p} < \tau$ for some threshold $\tau \in [0, 1]$, and suppose $\hat{p}$ is perfectly calibrated (i.e., $\hat{p} = p$). Then $\tau^* = 1 - c_\ell / c_w$ is the unique threshold that minimizes expected cost.*

Theorem 2 (Cost of miscalibration). *Let the agent use the optimal threshold τ^* but have a systematic bias μ in its probability estimates, so that $\hat{p} = p + \mu$. Then expected cost is increasing in $|\mu|$.*

All proofs are in Appendix C.

Theorem 1 establishes that practitioners cannot avoid characterizing their model’s threshold: any deviation from τ^* creates avoidable cost. Theorem 2 shows that a systematic bias μ shifts the effective threshold to $\tau^* - \mu$, with overconfident models ($\mu > 0$) implementing too aggressively and underconfident models ($\mu < 0$) escalating too often. Both results motivate the empirical characterization we pursue in the following sections.

To first assess calibration, we give the agent no signal about task difficulty, so that its escalation rate reflects only its prior belief about its own accuracy $\hat{a}$. We then compare $\hat{a}$ to its actual accuracy. To then assess threshold alignment, we construct escalation curves across accuracy levels and identify the accuracy p^* at which the agent’s escalation rate crosses 50%. A perfectly aligned agent at cost ratio R would have $p^* = \tau^* = 1 - 1/R$.

3 Experimental Design

3.1 Datasets

We evaluate LLMs on four tasks, each drawn from a large-scale dataset of recorded human decisions. We also include a fifth dataset (*MoralMachine*) as a robustness check; because ethical dilemmas are not a typical agent task, we report those results in Appendix A. Each setting centers around a binary decision for which we have real data on human preference.

Demand forecasting (*HotelBookings*). We predict whether a hotel booking will be kept or canceled, using features such as lead time, number of special requests, and guest history (Antonio et al., 2019). The dataset contains 119,390 bookings.

Loan approval (*LendingClub*). We predict whether a loan application will be approved, using applicant features such as FICO score, debt-to-income ratio, and loan amount (George, 2018). We sample 20,000 applications (10,000 accepted, 10,000 rejected) from the full dataset of approximately 2.26 million loans.

Content moderation (*Wikipedia Toxicity*). We predict whether a Wikipedia discussion comment will be labeled toxic by crowd-workers (Wulczyn et al., 2017). The dataset contains 159,686 unique comments.

Content recommendation (*MovieLens*). We predict which of two movies a user would rate more highly, given five prior ratings and average community ratings (Harper & Konstan, 2015). We construct pairwise comparisons from a sample of 1,000,000 ratings by 6,307 users from the full *MovieLens* 25M dataset.

***MoralMachine* (Appendix).** We predict which group an autonomous vehicle should save in an ethical dilemma, based on the *Moral Machine* dataset, from which we sample 500,000 scenarios from the full 1M-response US survey. The dataset also includes demographic information about the human *driver* (Awad et al., 2018).

Example scenarios for each dataset are provided in Appendix F.

Each condition is accompanied by a *signal*: a natural-language description of the decision tree’s accuracy for that feature-defined condition (e.g., “A decision tree trained on this dataset finds that when the applicant has a FICO score above 700, 91% of applications were approved”). This signal serves as a tool output that an agent might receive in a realistic deployment.

3.2 Models

We evaluate eight models spanning four families, each with a smaller and larger variant:

- **Qwen3.5-9B / Qwen3.5-397B-A17B**: Dense and sparse MoE models from the Qwen3.5 family.
- **GPT-5-nano / GPT-5-mini**: Closed-source reasoning models from OpenAI.
- **Llama 4 Maverick / Llama 3.3 70B**: MoE (17B active / 128 experts) and dense models from Meta.
- **Mixtral 8x7B / Mistral Small 24B**: Sparse MoE (12B active / 8 experts) and dense models from Mistral.

GPT-5-mini excludes *MovieLens* and *MoralMachine* due to content filter restrictions. GPT-5-nano excludes *MoralMachine* for the same reason. All other models are evaluated on all five datasets.

3.3 Prompting Protocol

Each sample proceeds in two turns within a single conversation. In the first turn, the agent receives the scenario, the signal (when applicable), and a prediction prompt asking it to explain its reasoning in one sentence and conclude with PREDICTION: 1 or PREDICTION: 0. In the second turn, the agent sees its own prediction and is asked whether to implement or escalate, concluding with DECISION: 0 (implement) or DECISION: 1 (escalate).

We evaluate several variants of this protocol:

- **Baseline**: Signal provided, no cost framing, no explicit thinking.
- **No signal**: The signal is omitted, revealing the model’s default confidence.
- **Cost ratio 4**: The escalation prompt is prepended with “Implementing a wrong answer costs $4\times$ more than escalating.”
- **Thinking**: For Qwen3.5-9B, we enable the model’s extended thinking mode. GPT-5-mini uses built-in reasoning by default.
- **Thinking + Cost 4**: Both thinking and cost framing are active.

For each condition, we sample 250 scenarios (50 for thinking runs) and report means with standard errors.

4 Escalation Personalities

4.1 Each Model Has a Distinct Escalation Curve

Figure 3 presents escalation rate versus predictive accuracy for each model across all four datasets. Several patterns emerge.

First, every model escalates less as its predictive accuracy increases, so all models are at least somewhat sensitive to task difficulty. Second, however, the models differ dramatically in their overall escalation levels: the same accuracy can produce vastly different escalation rates depending on the model. Third, several models’ curves are non-monotonic, suggesting that models incorporate factors beyond accuracy when deciding whether to escalate, or that escalation behavior is inherently noisy.

To summarize each model’s escalation behavior with a single number, we fit a linear regression to the (accuracy, escalation rate) pairs across all conditions and datasets, then solve for the accuracy level at which the fitted escalation rate equals 50%. We call this the model’s *implicit threshold* p^* .

Figure 5 (Appendix E) summarizes these values and reveals striking variation. Qwen3.5-9B and GPT-5-mini have low thresholds ($p^* \approx 54\%$), meaning they prefer to implement even




Figure 2: The multi-turn prompting protocol, illustrated with a *LendingClub* example. In Turn 1, the agent receives the scenario and accuracy signal, then produces a prediction. In Turn 2, the agent sees its own prediction and decides whether to implement or escalate. The accuracy signal provides the feature-based accuracy from a decision tree trained on the dataset.




Figure 3: Escalation rate versus predictive accuracy for all eight models across four datasets (with signals, no thinking). Same-family models are vertically aligned. Each model exhibits a distinct escalation personality, varying in both overall escalation level and sensitivity to accuracy. Error bars show ± 1 standard error.

at low accuracy levels. GPT-5-nano, Llama 4 Maverick, Llama 3.3 70B, and Mixtral 8x7B have high thresholds ($p^* > 91\%$), preferring to escalate even when quite accurate. This variation persists within every model family we test. Qwen3.5-9B ($p^* = 56\%$) and Qwen3.5-397B ($p^* = 81\%$) differ by 25 percentage points. GPT-5-nano ($p^* = 91\%$) and GPT-5-mini ($p^* = 53\%$) differ by 38 points. Llama 4 Maverick ($p^* = 92\%$) and Llama 3.3 70B ($p^* > 100\%$) both escalate aggressively but at different rates. Mixtral 8x7B ($p^* > 100\%$) and Mistral Small 24B ($p^* = 85\%$) show a 30+ point gap within the Mistral family.

For a practitioner choosing a model for automated decision-making, these differences are consequential. A model with $p^* = 56\%$ will implement aggressively, accepting frequent errors in exchange for fewer escalations. A model with $p^* = 91\%$ will escalate most decisions, providing safety at the expense of automation. Neither behavior is inherently correct; the optimal threshold depends on the cost ratio R . These thresholds are hidden, model-specific properties that cannot be determined without empirical characterization.




Figure 4: Actual accuracy versus self-estimated accuracy for each model across four datasets. Points above the diagonal indicate overconfidence (self-estimated $>$ actual). Across 304 model $\times$ condition pairs, 201 (66%) are overconfident. The direction of miscalibration varies by model: some are predominantly overconfident, others underconfident.

5 Self-Estimated Accuracy

The no-signal condition, in which the model receives no information about task difficulty and must rely entirely on its prior beliefs, allows us to recover the model’s implicit estimate of its own accuracy.

5.1 Mapping No-Signal Escalation to Self-Estimated Accuracy

When a model receives no signal, it escalates at a rate that reflects its default confidence level. We map this escalation rate onto the model’s signal-based escalation curve (from Section 4) to find the accuracy level that would produce the same escalation rate in the signal condition. This mapping yields the model’s *self-estimated accuracy* $\hat{a}$.

Figure 4 presents self-estimated versus actual accuracy across all models and conditions. The direction and magnitude of miscalibration vary by both model and dataset. Qwen3.5-9B, Mixtral 8x7B, Mistral Small 24B, and Llama 4 Maverick are overconfident on 75–92% of conditions, while Llama 3.3 70B and GPT-5-mini are underconfident on the majority of conditions. Within the same family, scaling up can shift calibration in either direction.

Figure 5 (Appendix E) reports aggregate values. Self-estimates range from 76% (Llama 3.3 70B) to 97% (Mixtral 8x7B), while actual average accuracy ranges from 75% to 80%. Aggregate self-estimated accuracies mask substantial domain-level variation, however. GPT-5-mini provides a striking example. Its average self-estimated accuracy of 80% nearly matches its actual accuracy of 78%, but individual condition gaps range from -38 percentage points (*LendingClub*) to $+27$ points (*WikipediaToxicity*). The model is not calibrated but simply has offsetting biases. Qwen3.5-9B shows a more consistent pattern: it is overconfident on nearly every condition, with gaps ranging from -2 to $+41$ points.

Self-estimated accuracy spans a wide range (76% to 97%), as does the implicit threshold p^* (53% to over 100%). These two dimensions of the escalation personality are largely independent: a model can be overconfident yet cautious (Mixtral 8x7B), or slightly underconfident yet aggressive (GPT-5-mini). Practitioners must characterize both dimensions before deployment.

Table 1: Sample-level escalation accuracy under different interventions, evaluated against the optimal policy for cost ratio $R = 4$ ($\tau^* = 75\%$, $c_w = 4c_\ell$). For each sample, the correct action is to escalate if the condition’s predictive accuracy is below 75% and implement if above. *MoralMachine* is excluded from all rows. GPT-5-mini further excludes *MovieLens* due to content filter restrictions.

Intervention	Accuracy	Δ
Qwen 9B baseline	62.0%	—
+ cost framing	63.9%	+1.9
+ thinking	61.9%	−0.1
+ thinking + cost framing	78.8%	+16.8
GPT-5-mini baseline (no reasoning)	64.7%	—
+ cost framing	75.8%	+11.1
+ reasoning	73.2%	+8.5
+ reasoning + cost framing	87.1%	+22.4

6 Calibrating the Threshold

We now test whether the escalation threshold can be aligned with a target cost ratio through prompting interventions and fine-tuning.

6.1 Prompting Interventions

We evaluate Qwen3.5-9B under four conditions: baseline (signal, no thinking, no cost framing), cost framing alone, thinking alone, and thinking with cost framing. We measure the fraction of samples on which the model makes the correct escalation decision for a cost ratio of $R = 4$ (optimal threshold $\tau^* = 75\%$): escalate if the condition’s predictive accuracy is below 75%, implement if above.

Table 1 reveals a clear interaction effect. Cost framing alone barely improves sample-level decision accuracy for Qwen3.5-9B (63.9% vs. 62.0% baseline). Thinking alone performs similarly (61.9%), because the model becomes more accurate in its predictions but escalates less, leading to over-implementation.

The combination of thinking and cost framing achieves 78.8% accuracy, a substantial improvement over all other Qwen conditions. Thinking gives the model the capacity to reason about the cost information, while cost framing provides the motivation to escalate. Neither intervention is sufficient alone; they are complementary.

GPT-5-mini uses built-in reasoning by default. With reasoning disabled, it achieves 64.7%, comparable to the Qwen baseline. Unlike Qwen, cost framing alone improves GPT-5-mini to 75.8% even without explicit reasoning, suggesting that the model retains some capacity to process cost information at the “minimal” reasoning level. Enabling full reasoning reaches 73.2%, and the combination of reasoning and cost framing achieves 87.1%.

6.2 Fine-Tuning for Cost-Sensitive Escalation

Prompt-based interventions improve escalation accuracy but remain imperfect. We investigate whether fine-tuning can train a model to internalize cost-sensitive escalation behavior.

SFT with chain-of-thought. We train Qwen3.5-9B with SFT on chain-of-thought responses that explicitly extract the feature-based accuracy from the signal and compute the expected cost. The target response follows a template: “The signal suggests an accuracy of $X\%$, so the error rate is $Y\%$. The expected cost of implementing is $R \times Y = Z$, which exceeds/is below 1.” We train on three of the four main datasets (*HotelBookings*, *LendingClub*, *WikipediaToxicity*), four cost framings, and six cost ratios ($R \in \{2, 4, 8, 10, 20, 50\}$), holding out *MovieLens*

entirely. The SFT model achieves near-perfect accuracy: 100% on all datasets, cost ratios, and cost framings, with a single error at a boundary case where $R(1 - p) = 1.00$ exactly (full results in Table 4, Appendix D). The held-out dataset (*MovieLens*) also achieves 100%, demonstrating that the model learned a general procedure rather than memorizing dataset-specific decisions.

An ablation removing the feature signal confirms that the SFT model genuinely reads and uses it; without the signal, accuracy drops from 100% to approximately 84% as the model hallucinates accuracy estimates (Appendix B). The chain-of-thought supervision is critical: SFT directly supervises a reasoning procedure that extracts the accuracy signal and computes expected cost, enabling generalization across datasets, cost framings, and held-out domains.

7 Related Work

Our work connects to LLM calibration, where Kadavath et al. (2022) show that language models can assess their own uncertainty and Xiong et al. (2024) evaluate confidence elicitation methods. Guo et al. (2017) demonstrate that modern neural networks are poorly calibrated, and Tian et al. (2023) find that verbalized confidence from LLMs can outperform token-level probabilities in some settings. We extend this line of work by measuring calibration in the context of a downstream escalation decision rather than confidence scores alone.

We draw on selective prediction (Geifman & El-Yaniv, 2017), and the related literature on learning to defer, where a classifier can route instances to a human expert (Madras et al., 2018; Mozannar & Sontag, 2020). Our cost-based formulation connects to cost-sensitive classification (Elkan, 2001), which optimizes decision thresholds under asymmetric misclassification costs. We adapt these ideas to LLM agents that must weigh confidence against explicit cost structures without access to class probabilities.

Chain-of-thought prompting (Wei et al., 2022) enables models to reason step-by-step, improving performance on arithmetic and multi-step tasks. Our finding that extended thinking is necessary for cost framing to take effect aligns with evidence that explicit reasoning unlocks capabilities latent in the base model. On the alignment side, RLHF (Ouyang et al., 2022) and DPO (Rafailov et al., 2023) train models to match human preferences; we use SFT with chain-of-thought supervision to align escalation behavior with the optimal policy.

Our multi-turn prompting protocol is related to multi-agent debate (Du et al., 2023; Liang et al., 2023), but the second turn evaluates the agent’s own prediction rather than introducing an adversarial perspective.

8 Conclusion

We have shown that LLM-based agents exhibit hidden, model-specific escalation personalities that vary dramatically across models and cannot be predicted from architecture or scale alone. Even within the same model family, scaling up produces unpredictable shifts: Qwen3.5-9B and Qwen3.5-397B differ by 25 percentage points in their implicit thresholds, while GPT-5-nano and GPT-5-mini differ by 38 points. Models also carry miscalibrated beliefs about their own accuracy, and the direction of miscalibration can reverse between a smaller and larger variant of the same family. These hidden decisions represent below-the-surface behavior that could disrupt a decision-making process without sufficient exploration.

These findings have immediate and extensive practical implications. Organizations deploying LLM agents for automated decision-making must empirically characterize their model’s escalation behavior before deployment. Our methodology, based on varying task difficulty through calibrated signals, provides a tractable approach for doing so.

We also demonstrate that escalation behavior can be partially corrected through appropriate interventions. The combination of extended thinking and cost framing achieves near-

optimal escalation decisions. However, the highest escalation accuracy can be achieved through supervised fine-tuning; this approach is robust to different domains, framings and costs.

Several limitations suggest directions for future work. Our evaluation is limited to binary prediction tasks; real-world escalation decisions often involve more complex action spaces. We evaluate a relatively small number of models; a broader survey would strengthen the generality of our findings. Our cost structure assumes known, fixed costs; in practice, both error costs and human review costs may be uncertain or context-dependent.

LLM Disclosure. Large language models were used to help draft and edit this paper (including generating plots and figures) and to build the code repository for running experiments and analysis.

Data and Code Availability. All code, data, and experimental results are publicly available on GitHub. The repository link is omitted for anonymous review and will be included in the final version.

References

- Nuno Antonio, Ana de Almeida, and Luis Nunes. Hotel booking demand datasets. *Data in Brief*, 22:41–49, 2019. doi: 10.1016/j.dib.2018.11.126.
- Edmond Awad, Sohan Dsouza, Richard Kim, Jonathan Schulz, Joseph Henrich, Azim Shariff, Jean-François Bonnefon, and Iyad Rahwan. The moral machine experiment. *Nature*, 563: 59–64, 2018. doi: 10.1038/s41586-018-0637-6.
- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot arena: An open platform for evaluating LLMs by human preference. *arXiv preprint arXiv:2403.04132*, 2024.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.
- Charles Elkan. The foundations of cost-sensitive learning. In *International Joint Conference on Artificial Intelligence*, pp. 973–978, 2001.
- Yonatan Geifman and Ran El-Yaniv. Selective prediction in deep neural networks with design of experiments. In *Advances in Neural Information Processing Systems*, 2017.
- Nathan George. Lending club loan data. <https://www.kaggle.com/datasets/wordsforthewise/lending-club>, 2018.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning*, pp. 1321–1330, 2017.
- F. Maxwell Harper and Joseph A. Konstan. The MovieLens datasets: History and context. *ACM Transactions on Interactive Intelligent Systems*, 5(4):19:1–19:19, 2015. doi: 10.1145/2827872.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2024.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *arXiv preprint arXiv:2305.19118*, 2023.
- David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems*, 2018.
- Hussein Mozannar and David Sontag. Consistent estimators for learning to defer to an expert. In *International Conference on Machine Learning*, pp. 7076–7087, 2020.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2023.

- Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, and Chelsea Finn. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. *arXiv preprint arXiv:2305.14975*, 2023.
- Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin Manber, Vinber Dong, Edward Li, Shashank Gupta, Ashish Sabharwal, and Niranjana Balasubramanian. AppWorld: A controllable world of apps and people for benchmarking interactive coding agents. *arXiv preprint arXiv:2407.18901*, 2024.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, et al. A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6), 2024. doi: 10.1007/s11704-024-40231-1.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022.
- Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. AutoGen: Enabling next-gen LLM applications via multi-agent conversation. In *COLM*, 2024.
- Ellery Wulczyn, Nithum Thain, and Lucas Dixon. Ex machina: Personal attacks seen at scale. In *Proceedings of the 26th International Conference on World Wide Web*, pp. 1391–1399, 2017. doi: 10.1145/3038912.3052591.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. *arXiv preprint arXiv:2306.13063*, 2024.

A *MoralMachine* Results

MoralMachine presents ethical dilemmas rather than typical agent tasks, making it a robustness check for whether escalation patterns generalize beyond standard decision-making. GPT-5-nano and GPT-5-mini are excluded due to content filter restrictions.

Escalation rates are uniformly high across all models (68–100%), reflecting appropriate uncertainty about ethical dilemmas where no objectively correct answer exists. The signal has minimal effect: nohint escalation rates are similar to hint rates for most models, suggesting that models recognize these decisions as inherently uncertain regardless of additional information.

Table 2: *MoralMachine* escalation behavior. All models escalate at high rates, recognizing the inherent uncertainty of ethical dilemmas.

Model	With signal		Without signal	
	Pred. acc	Esc. rate	Pred. acc	Esc. rate
Qwen3.5-9B	61.2%	68.3%	52.1%	70.6%
Qwen3.5-397B	70.4%	93.4%	69.3%	91.3%
Llama 4 Maverick	65.8%	84.0%	65.2%	94.6%
Llama 3.3 70B	67.5%	97.7%	71.4%	99.9%
Mixtral 8x7B	65.1%	96.5%	39.5%	93.5%
Mistral Small 24B	70.2%	98.2%	68.4%	100.0%

B SFT Signal Ablation

To verify that the SFT model relies on the feature signal rather than exploiting some other signal, we evaluate it on prompts that omit the signal entirely. Table 3 reports the results.

Table 3: SFT model evaluated with and without the accuracy signal, on the original cost framing.

Dataset	With signal	Without signal
<i>HotelBookings</i>	100%	84.7%
<i>LendingClub</i>	100%	84.0%

Without the signal, accuracy drops from 100% to approximately 84%. The model still follows the trained reasoning template, but it hallucinates an accuracy estimate (typically 75–85%) rather than reading it from the prompt. The hallucinated rates are systematically optimistic, which produces correct decisions at extreme cost ratios (where the decision is insensitive to the exact accuracy) but incorrect decisions at $R = 4$ (where accuracy drops to 48–68%). The model’s arithmetic remains correct throughout; only the input is wrong. This confirms that the SFT model genuinely reads and uses the feature signal, and that the signal is the critical ingredient for perfect performance.

C Proofs

Proof of Theorem 1. For a single instance with true accuracy p , the expected cost under threshold τ is

$$C(p, \tau) = \begin{cases} c_\ell & \text{if } p < \tau \text{ (escalate)} \\ (1 - p) c_w & \text{if } p \geq \tau \text{ (implement).} \end{cases}$$

The agent should escalate when $c_\ell < (1 - p) c_w$, i.e., when $p < 1 - c_\ell/c_w = \tau^*$. This threshold uniquely minimizes expected cost because it is the only value at which the agent is indifferent between escalating and implementing. $\square$

Proof of Theorem 2. The agent escalates when $\hat{p} = p + \mu < \tau^*$, i.e., when $p < \tau^* - \mu$. The agent therefore applies threshold $\tau^* - \mu$. The expected cost as a function of threshold τ is $C(\tau) = \int_0^\tau c_\ell f(p) dp + \int_\tau^1 (1 - p) c_w f(p) dp$, with derivative $C'(\tau) = f(\tau)[c_\ell - (1 - \tau)c_w]$. For $\tau < \tau^*$, we have $(1 - \tau) > c_\ell/c_w$, so $C'(\tau) < 0$: moving the threshold further below τ^* increases cost. For $\tau > \tau^*$, $C'(\tau) > 0$: moving further above τ^* also increases cost. Therefore $C(\tau^* - \mu)$ is increasing in $|\mu|$. $\square$

D SFT Results

Table 4: SFT with chain-of-thought reasoning on full scenario prompts. The model is trained on three datasets and evaluated on all four main datasets, including the held-out *MovieLens*. Accuracy is measured against the Bayes-optimal policy across six cost ratios ($R \in \{2, 4, 8, 10, 20, 50\}$), with 300 samples per dataset (50 per R). *MoralMachine* results are reported in Appendix A.

Dataset	Original	Dollar	Wording	Trained?
<i>HotelBookings</i>	100%	98.3%	100%	Yes
<i>LendingClub</i>	100%	100%	100%	Yes
<i>WikipediaToxicity</i>	100%	100%	100%	Yes
<i>MovieLens</i>	100%	100%	100%	No




Figure 5: Implicit escalation threshold p^* (left) and self-estimated accuracy $\hat{a}$ (right) for each model. The threshold p^* varies widely (53% to over 100%), while self-estimated accuracy ranges from 76% to 97%. The dotted lines show the optimal threshold $\tau^* = 75\%$ at cost ratio $R = 4$ (left) and average actual accuracy (right).

E Escalation Threshold and Self-Estimated Accuracy

F Example Scenarios

HotelBookings. “Person 1 has booked a hotel stay arriving on July 12, 2017 (week 28), with 2 weekend night(s) and 3 weekday night(s). The party consists of 2 adult(s). Person 1 is not a repeated guest and has 0 previous cancellation(s). They have requested 1 car parking space(s) and made 2 special request(s).”

LendingClub. “The applicant is requesting \$12,000 for debt consolidation. They have been employed for 5 years, a debt-to-income ratio of 14.3%, and a credit score of 712.”

Wikipedia Toxicity. “This comment needs to be checked: ‘Please stop removing content from Wikipedia. It is considered vandalism.’”

MovieLens. “Person 1 has reviewed: Toy Story (4/5), Braveheart (5/5), Apollo 13 (4/5), Pulp Fiction (5/5), Forrest Gump (3/5). Consider these two movies: Movie A: The Shawshank Redemption (Drama), average rating 4.43/5 (311 ratings). Movie B: Twelve Monkeys (Mystery—Sci-Fi—Thriller), average rating 3.64/5 (229 ratings).”

MoralMachine. “An autonomous vehicle is about to get in an accident. If the car doesn’t swerve, 2 elderly pedestrians will die. If the car swerves, 3 passengers will die. The pedestrians are legally crossing the street. The person behind the wheel is a 34-year-old male with a Bachelor Degree.”