

# Act or Escalate? Evaluating Escalation Behavior in Automation with Language Models

Anonymous authors

Paper under double-blind review

## Abstract

Effective automation hinges on deciding when to act and when to escalate. We model this as a decision under uncertainty: an LLM forms a prediction, estimates its probability of being correct, and compares the expected costs of acting and escalating. Using this framework across five domains of recorded human decisions—demand forecasting, content recommendation, content moderation, loan approval, and autonomous driving—and across multiple model families, we find marked differences in the implicit thresholds models use to trade off these costs. These thresholds vary substantially and are not predicted by architecture or scale, while self-estimates range from well-calibrated to systematically miscalibrated in opposite directions across models. We then test interventions that target this decision process by varying cost ratios, providing accuracy signals, and training models to follow the desired escalation rule. Prompting helps mainly for reasoning models. SFT on chain-of-thought targets yields the most robust policies, which generalize across datasets, cost ratios, prompt framings, and held-out domains. These results suggest that escalation behavior is a model-specific property that should be characterized before deployment, and that robust alignment benefits from training models to reason explicitly about uncertainty and decision costs.

## 1 Introduction

Agents based on large language models (LLMs) are increasingly deployed to automate consequential decisions, from performing autonomous tasks (Wang et al., 2024) to generating code (Jimenez et al., 2024) and managing enterprise workflows (Wu et al., 2024). Most evaluations focus on speed, accuracy, and cost-savings (Jimenez et al., 2024; Chiang et al., 2024; Trivedi et al., 2024). However, in each setting, an agent faces a choice that receives far less scrutiny: should it implement its own decision, or escalate to a human?

Agent escalation behavior is fundamental to effective automation. An agent that fails to escalate when it is uncertain or incorrect will propagate errors at scale (Lanham et al., 2023), while an agent that always escalates never reduces the human workload it was meant to replace. Two parameters govern effective escalation. First, the agent must have a calibrated sense of its own accuracy, recognizing when its predictions are likely wrong and escalation is appropriate. Second, the agent must weigh the cost of implementing an error against the cost of escalating to a human, and defer when the risk outweighs the savings.

In the main text we report results for eight models in four families, each represented by a smaller and larger variant: Qwen3.5-4B and Qwen3.5-9B, Gemma 3 4B and Gemma 3 12B, Claude Sonnet 4.6 and Claude Opus 4.7, and GPT-5-nano and GPT-5-mini. The Appendix extends the analysis to five additional models: the larger MoE flagship Qwen3.5-397B-A17B, plus Llama 4 Maverick and Llama 3.3 70B, and Mixtral 8x7B and Mistral Small 24B. We evaluate these models on five decision-making tasks, each derived from large-scale human decision data: demand forecasting, loan approval, content moderation, content recommendation, and moral dilemmas (featured in the Appendix). We find that LLMs are both miscalibrated and inconsistent in their escalation behavior.




Figure 1: The escalation decision. A human delegates a task to an LLM agent, which produces a prediction and then decides whether to implement (incurring error cost  $c_w$  if wrong) or escalate (incurring labor cost  $c_\ell$ ). The agent escalates when its estimated probability of being correct  $\hat{p}$  falls below a threshold  $\tau$ .

First, models differ widely in self-assessment of their own predictive accuracy. Some systematically overestimate it, others underestimate it, and others are well-calibrated. The direction and magnitude of miscalibration vary by model and domain. Prior work has documented LLM overconfidence in factual knowledge (Kadavath et al., 2022) and calibration failures in question answering (Xiong et al., 2024), but these studies measure confidence in isolation. We show that miscalibration has direct operational consequences: overconfident models implement predictions they should defer, while underconfident models escalate decisions they could handle.

Second, escalation behavior varies substantially across different models. This variance is not associated with model architecture or size. Some models tend to prefer escalating (lower decision threshold), some generally favor implementing (higher decision threshold). This variation persists even within model families. Scaling up does not consistently shift the threshold in either direction. Together, this miscalibration and variance in escalation behavior constitute latent model dynamics that could disrupt a larger workflow.

These dynamics, however, can be corrected. We test a series of interventions that target the escalation decision. Prompt-based cost framing alone has no effect, but combining it with extended thinking produces improvement for reasoning models. Supervised fine-tuning (SFT) on chain-of-thought targets proves most effective: the resulting model makes near-optimal escalation decisions across all datasets, cost ratios, and prompt framings, and generalizes to held-out domains it was never trained on. Together, our results establish that escalation behavior is a model-specific property that should be characterized before deployment, and that aligning it robustly requires training models to reason explicitly about uncertainty and decision costs.

## 2 Theory

### 2.1 The escalation decision

Consider a workflow in which a human delegates a binary task to an LLM-based agent. The agent observes a scenario  $x$  and produces a prediction  $\hat{y} \in \{0, 1\}$ . The ground truth is the human’s preference  $y \in \{0, 1\}$ . The agent then decides whether to *implement* (act on) its prediction or *escalate* (defer to the human). Every action an agent takes in practice contains an implicit escalation decision: when it makes a tool call, generates a response, or commits a change, it is choosing to implement rather than escalate and ask for guidance. Figure 1 illustrates this workflow.

We model the agent’s behavior as follows. After producing its prediction, the agent forms an internal estimate  $\hat{p}$  of its probability of being correct  $P(\hat{y} = y)$ . The cost structure then determines an optimal decision threshold  $\tau$  (derived below): the agent escalates if  $\hat{p} < \tau$  and implements otherwise.

## 2.2 Cost structure

One of two costs can be incurred depending on the agent’s action. If the agent escalates, the human pays a labor cost  $c_\ell > 0$ , which represents the time and effort required of the human. If the agent implements and its prediction is wrong, the human pays an error cost  $c_w > c_\ell$ . Correct implementations incur zero cost.

**Theorem 1** (Uniqueness of the optimal threshold). *Let the agent escalate whenever  $\hat{p} < \tau$  for some threshold  $\tau \in [0, 1]$ , and suppose  $\hat{p}$  is perfectly calibrated (i.e.,  $\hat{p} = p$ ). Then  $\tau^* = 1 - c_\ell / c_w$  is the unique threshold that minimizes expected cost.*

**Theorem 2** (Cost of miscalibration). *Let the agent use the optimal threshold  $\tau^*$  but have a systematic bias  $\mu$  in its probability estimates, so that  $\hat{p} = p + \mu$ . Then expected cost is increasing in  $|\mu|$ .*

All proofs are in Appendix C.

Theorem 1 establishes that practitioners cannot avoid characterizing their model’s threshold: any deviation from  $\tau^*$  creates avoidable cost. Theorem 2 shows that a systematic bias  $\mu$  shifts the effective threshold to  $\tau^* - \mu$ , with overconfident models ( $\mu > 0$ ) implementing too aggressively and underconfident models ( $\mu < 0$ ) escalating too often. Both results motivate the empirical characterization we pursue in the following sections.

To first assess calibration, we give the agent no signal, so that its escalation rate reflects only its prior belief about its own accuracy  $\hat{a}$ . We then compare  $\hat{a}$  to its actual accuracy. To then assess threshold alignment, we construct escalation curves across accuracy levels and identify the accuracy  $p^*$  at which the agent’s escalation rate crosses 50%. A perfectly aligned agent at cost ratio  $R$  would have  $p^* = \tau^* = 1 - 1/R$ .

## 3 Experimental design

### 3.1 Datasets

We evaluate LLMs on four tasks, each drawn from a large-scale dataset of recorded human decisions. We also include a fifth dataset (*MoralMachine*) as a robustness check; because ethical dilemmas are not a typical agent task, we report those results in Appendix A. Each setting centers around a binary decision for which we have real data on human preference.

**Demand forecasting (*HotelBookings*).** We predict whether a hotel booking will be kept or canceled, using features such as lead time, number of special requests, and guest history (Antonio et al., 2019). The dataset contains 119,390 bookings.

**Loan approval (*LendingClub*).** We predict whether a loan application will be approved, using applicant features such as FICO score, debt-to-income ratio, and loan amount (George, 2018). We sample 20,000 applications (10,000 accepted, 10,000 rejected) from the full dataset of approximately 2.26 million loans.

**Content moderation (*Wikipedia Toxicity*).** We predict whether a Wikipedia discussion comment will be labeled toxic by crowd-workers (Wulczyn et al., 2017). The dataset contains 159,686 unique comments.

**Content recommendation (*MovieLens*).** We predict which of two movies a user would rate more highly, given five prior ratings and average community ratings (Harper & Konstan, 2015). We construct pairwise comparisons from a sample of 1,000,000 ratings by 6,307 users from the full *MovieLens* 25M dataset.

***MoralMachine* (Appendix).** We predict which group an autonomous vehicle should save in an ethical dilemma, based on the *Moral Machine* dataset, from which we sample 500,000 scenarios from the full 1M-response US survey. The dataset also includes demographic information about the human *driver* (Awad et al., 2018).

Example scenarios for each dataset are provided in Appendix G.

### 3.2 Models

The main text reports eight models in four families, each with a smaller and larger variant:

- **Qwen3.5-4B / Qwen3.5-9B**: Dense models from the Qwen3.5 family.
- **Gemma 3 4B / Gemma 3 12B**: Dense models from Google.
- **Claude Sonnet 4.6 / Claude Opus 4.7**: Closed-source models from Anthropic.
- **GPT-5-nano / GPT-5-mini**: Closed-source reasoning models from OpenAI.

The Appendix extends the analysis to the MoE flagship Qwen3.5-397B-A17B and two more families: Llama 4 Maverick / Llama 3.3 70B (Meta) and Mixtral 8x7B / Mistral Small 24B (Mistral).

GPT-5-mini excludes *MovieLens* and *MoralMachine* due to content filter restrictions. GPT-5-nano excludes *MoralMachine* for the same reason. Gemma 3 and Claude models are evaluated on the four main datasets but not *MoralMachine*. The remaining models are evaluated on all five datasets.

### 3.3 Prompting protocol

Each scenario is accompanied by a *signal*: a natural-language summary of a decision tree’s output for the scenario’s feature-defined condition, including the feature split and its predictive accuracy (e.g., “A decision tree trained on this dataset finds that when the applicant has a FICO score above 700, 91% of applications were approved”). This signal allows us to isolate escalation behavior from the model’s own beliefs about its accuracy by providing an external reference point. It also reflects what an agent might receive in a realistic deployment, where it could receive signals from the output of a tool that assists in a task.

Each sample proceeds in two turns within a single conversation. In the first turn, the agent receives the scenario, the signal (when applicable), and a prediction prompt asking it to explain its reasoning in one sentence and conclude with PREDICTION: 1 or PREDICTION: 0. In the second turn, the agent sees its own prediction and is asked whether to implement or escalate, concluding with DECISION: 0 (implement) or DECISION: 1 (escalate).

We evaluate several variants of this protocol. **Baseline** and **No signal** are used in Sections 4 and 5 to characterize each model’s intrinsic escalation behavior. **Cost framing**, **Thinking**, and their combination are interventions evaluated in Section 6.

- **Baseline**: Signal provided, no cost framing, no explicit thinking.
- **No signal**: The signal is omitted, revealing the model’s default confidence.
- **Cost framing** ( $R = 4$ ): The escalation prompt is prepended with “Implementing a wrong answer costs  $4\times$  more than escalating.”
- **Thinking**: For Qwen3.5-9B, we enable the model’s extended thinking mode. For GPT-5-mini, we increase reasoning\_effort from minimal to medium.
- **Thinking + cost framing**: Both interventions are active.

For each condition, we sample 250 scenarios (100 for thinking runs) and report means with standard errors.

### 3.4 Metrics

We use three accuracy measures throughout, and refer to each by its specific name to avoid ambiguity.




Figure 2: The multi-turn prompting protocol, illustrated with a *LendingClub* example. In Turn 1, the agent receives the scenario and signal, then produces a prediction. In Turn 2, the agent sees its own prediction and decides whether to implement or escalate. The signal provides the feature split and predictive accuracy from a decision tree trained on the dataset.

**Predictive accuracy** ( $p$ ) is the rate at which the agent's prediction  $\hat{y}$  matches the ground truth  $y$  within a condition. It measures how well the agent's prediction tracks the underlying labels and depends only on the prediction in Turn 1.

**Escalation accuracy** is the rate at which the agent's implement/escalate choice in Turn 2 agrees with the cost-optimal action for a given cost ratio  $R$ . An escalation is optimal when the prediction's expected error cost exceeds the labor cost, i.e., when  $p < \tau^* = 1 - 1/R$ . Escalation accuracy is therefore a property of the joint behavior across both turns, evaluated against the Bayes-optimal policy.

**Self-estimated predictive accuracy** ( $\hat{a}$ ) is the agent's implicit belief about its own predictive accuracy, recovered by mapping its no-signal escalation rate onto the signal-based escalation curve (Section 5).

## 4 Escalation profiles

Models form internal beliefs about their own predictive accuracy, and these beliefs shape their escalation decisions. To isolate escalation behavior from these beliefs, we provide an explicit signal that specifies the predictive accuracy for each scenario. Figure 3 presents the resulting escalation rate versus predictive accuracy for each model across all four datasets. Several patterns emerge.

First, every model escalates less as its predictive accuracy increases, so all models are at least somewhat sensitive to the signal. Second, however, the models differ dramatically in their overall escalation levels: the same accuracy can produce vastly different escalation rates depending on the model. Third, several models' curves are non-monotonic, suggesting that models incorporate factors beyond accuracy when deciding whether to escalate, or that escalation behavior is inherently noisy.

To summarize each model's escalation behavior with a single number, we fit a linear regression to the (accuracy, escalation rate) pairs across all conditions and datasets, then solve for the accuracy level at which the fitted escalation rate equals 50%. We call this the model's *implicit threshold*  $p^*$ .

Figure 6 (Appendix E) summarizes these values and reveals striking variation. Across our eight main-text models,  $p^*$  ranges from 53% (GPT-5-mini) to 98% (Gemma 3 12B), spanning nearly the full plausible range. Scaling within a family shifts behavior unpredictably, and not only in threshold magnitude. Qwen3.5-4B ( $p^* = 61\%$ ) and Qwen3.5-9B ( $p^* = 56\%$ ) sit close in threshold but differ sharply in the *shape* of their escalation curves: 4B approximates a step function, escalating almost always below its threshold and almost never above, while 9B responds gradually and non-monotonically. Gemma 3 shows a large threshold gap: 4B sits at  $p^* = 76\%$  while 12B jumps to  $p^* = 98\%$ . GPT-5-nano ( $p^* = 91\%$ ) and GPT-5-mini ( $p^* = 53\%$ ) differ by 38 points. The Claude family is more tightly clustered, with Sonnet at






Figure 3: Each model exhibits a distinct escalation profile. Escalation rate versus predictive accuracy for all eight models across four datasets (with signals, no thinking). Same-family models are vertically aligned. Models vary in both overall escalation level and sensitivity to the signal. Error bars show  $\pm 1$  standard error.

204  $p^* = 66\%$  and Opus at  $p^* = 59\%$ , both leaning toward implementation. Appendix F reports  
 205 the same analysis for five additional models, including the Qwen3.5-397B MoE flagship  
 206 ( $p^* = 81\%$ , a 25-point gap above the dense 9B) and the Llama and Mistral families, where  
 207 within-family gaps of 30 to 40 percentage points also recur.

208 For a practitioner choosing a model for automated decision-making, these differences are  
 209 consequential. A model with  $p^* = 56\%$  will implement aggressively, accepting frequent  
 210 errors in exchange for fewer escalations. A model with  $p^* = 91\%$  will escalate most  
 211 decisions, providing safety at the expense of automation. Neither behavior is inherently  
 212 correct; the optimal threshold depends on the cost ratio  $R$ . These thresholds are latent,  
 213 model-specific properties that cannot be determined without empirical characterization.

## 214 5 Self-estimated predictive accuracy

215 Having established signal-based escalation curves, we now recover what a model believes  
 216 about its own predictive accuracy. Without a signal, the model escalates at a rate that reflects  
 217 only its prior beliefs. We map this no-signal escalation rate onto the signal-based curve  
 218 to find the predictive accuracy that would produce the same escalation rate, yielding the  
 219 model’s *self-estimated predictive accuracy*  $\hat{a}$ .

220 We find that calibration varies markedly across our main-text models (Figure 4). Qwen3.5-9B  
 221 is the most overconfident, with a 19-point gap between its self-estimated and actual predic-  
 222 tive accuracy (94% vs. 75%). Its smaller sibling Qwen3.5-4B is only mildly overconfident  
 223 (82% vs. 77%), so scaling the dense Qwen up from 4B to 9B sharply worsens calibration.  
 224 Both Gemma 3 sizes are mildly underconfident (self-estimates of 68–69% against actual  
 225 75–76%). GPT-5-nano is moderately overconfident (85% vs. 76%). Claude Sonnet 4.6, Claude  
 226 Opus 4.7, and GPT-5-mini sit close to calibrated, with self-estimates within 3 percentage  
 227 points of actual. Within the same family, scaling up shifts calibration unpredictably: dense  
 228 Qwen worsens, GPT-5 moves from overconfident to calibrated, and Gemma 3 stays un-  
 229 derconfident at both sizes. Appendix F extends this analysis to five additional models,  
 230 including Qwen3.5-397B and the Llama and Mistral families.

231 These aggregate numbers, however, understate the problem (Figure 6, Appendix E). Self-  
 232 estimates span 68% (Gemma 3 4B) to 94% (Qwen3.5-9B), while actual average predictive  
 233 accuracy stays narrowly within 75% to 80%. Qwen3.5-9B exemplifies severe overconfidence:  
 234 it is overconfident on nearly every condition, with gaps ranging from  $-2$  to  $+41$  percentage




Figure 4: Actual versus self-estimated predictive accuracy for each of the eight main-text models across four datasets. Points above the diagonal indicate overconfidence (self-estimated > actual). The direction and magnitude of miscalibration vary markedly by model: Qwen3.5-9B is consistently overconfident, Gemma 3 is underconfident, GPT-5-nano is moderately overconfident, and Claude and GPT-5-mini are close to calibrated. Appendix F shows the same plot for the additional four models.

points. Aggregate calibration can also mask per-condition variation: Claude Sonnet appears nearly calibrated overall, yet its condition-level gaps still span both directions. Appendix F discusses the additional models, including cases where offsetting biases hide miscalibration in the aggregate.

Even knowing a model’s self-estimated predictive accuracy does not predict its escalation threshold. Self-estimated predictive accuracy spans 68% to 94%, and  $p^*$  spans 53% to 98%, but the two are largely independent. Gemma 3 12B is underconfident yet has the highest threshold (escalates aggressively), while GPT-5-mini is nearly calibrated yet has the lowest threshold (implements aggressively). Practitioners must characterize both dimensions before deployment.

## 6 Calibrating the threshold

We now test whether the escalation threshold can be aligned with a target cost ratio through prompting interventions and fine-tuning.

### 6.1 Prompting interventions

We evaluate Qwen3.5-9B under four conditions: baseline (signal, no thinking, no cost framing), cost framing alone, thinking alone, and thinking with cost framing. For each sample, we score whether the model makes the cost-optimal escalation decision at  $R = 4$ , which sets the optimal threshold at  $\tau^* = 75\%$ .

Table 1 reveals a clear interaction effect. Cost framing alone barely improves escalation accuracy for Qwen3.5-9B (59.8% vs. 59.0% baseline). Thinking alone is slightly worse than baseline (58.1%), because the model’s predictive accuracy rises but its escalation rate drops, leading to over-implementation.

The combination of thinking and cost framing achieves 80.5% escalation accuracy, a substantial improvement over all other Qwen conditions. Thinking gives the model the capacity to reason about the cost information, while cost framing provides the motivation to escalate. Neither intervention is sufficient alone; they are complementary.




Figure 5: Cost framing intervention. The cost statement (red box) is prepended to the Turn 2 user message, exposing the cost ratio  $R = 4$  to the agent. All other prompt elements are unchanged from the baseline protocol (Figure 2).

Table 1: Sample-level escalation accuracy under different interventions, evaluated against the optimal policy for cost ratio  $R = 4$  ( $\tau^* = 75\%$ ,  $c_w = 4c_\ell$ ). For each sample, the correct action is to escalate if the condition's predictive accuracy is below 75% and implement if above. *MoralMachine* is excluded from all rows. GPT-5-mini further excludes *MovieLens* due to content filter restrictions.

| Intervention               | Accuracy (95% CI)  | $\Delta$ |
|----------------------------|--------------------|----------|
| Qwen 9B baseline           | 59.0% [58.3, 59.7] | —        |
| + cost framing             | 59.8% [59.1, 60.5] | +0.8     |
| + thinking                 | 58.1% [56.7, 59.4] | −0.9     |
| + thinking + cost framing  | 80.5% [79.4, 81.5] | +21.5    |
| GPT-5-mini baseline        | 66.3% [65.3, 67.4] | —        |
| + cost framing             | 77.7% [76.7, 78.6] | +11.4    |
| + reasoning                | 71.2% [70.1, 72.3] | +4.9     |
| + reasoning + cost framing | 91.4% [90.4, 92.3] | +25.1    |

GPT-5-mini uses built-in reasoning controlled by OpenAI's reasoning\_effort parameter. The lowest available setting is minimal (OpenAI does not support disabling reasoning entirely), so some residual reasoning may contribute to the baseline. At the minimal level, it achieves 66.3% escalation accuracy, comparable to the Qwen baseline. Unlike Qwen, cost framing alone improves GPT-5-mini to 77.7% even at minimal reasoning, suggesting that the model retains some capacity to process cost information without extended reasoning. At medium reasoning, escalation accuracy reaches 71.2%, and the combination of medium reasoning and cost framing achieves 91.4%.

## 6.2 Fine-tuning for cost-sensitive escalation

Prompt-based interventions improve escalation accuracy but remain imperfect. We investigate whether fine-tuning can train a model to internalize cost-sensitive escalation behavior. We train three SFT models with chain-of-thought responses that explicitly extract the predictive accuracy from the signal and compute the expected cost: Qwen2.5-7B-Instruct, Qwen3.5-4B, and Qwen3.5-9B. The target response follows a template: "The signal suggests a predictive accuracy of X%, so the error rate is Y%. The expected cost of implementing is  $R \times Y = Z$ , which exceeds/is below 1." We train on three of the four main datasets (*HotelBookings*, *LendingClub*, *WikipediaToxicity*), four cost framings, and six cost ratios ( $R \in \{2, 4, 8, 10, 20, 50\}$ ), holding out *MovieLens* entirely.

All three SFT models substantially outperform the prompt-only interventions in Table 1. Qwen2.5-7B-Instruct achieves near-perfect escalation accuracy: 100% on most cells with a single error at a boundary case where  $R(1 - p) = 1.00$  exactly. Qwen3.5-4B reaches 95.7%



overall (held-out *MovieLens* 96.7%), and Qwen3.5-9B reaches 87.3% overall (held-out 89.2%). The held-out generalization across all three models indicates that SFT teaches a general procedure rather than memorizing dataset-specific decisions. The within-Qwen3.5 ranking is non-monotonic, with 4B exceeding 9B, suggesting the same recipe may benefit from per-model tuning. Full per-dataset and per-framing results appear in Table 4, Appendix D.

To confirm that the SFT model reads and uses the signal, we run an ablation that removes it entirely. Without the signal, escalation accuracy drops from 100% to approximately 84% as the model hallucinates predictive-accuracy estimates instead of reading them from the signal (Appendix B). The chain-of-thought supervision is critical: SFT directly supervises a reasoning procedure that extracts the predictive accuracy from the signal and computes expected cost, enabling generalization across datasets, cost framings, and held-out domains.

## 7 Related work

Our work connects to LLM calibration, where Kadavath et al. (2022) show that language models can assess their own uncertainty and Xiong et al. (2024) evaluate confidence elicitation methods. Guo et al. (2017) demonstrate that modern neural networks are poorly calibrated, and Tian et al. (2023) find that verbalized confidence from LLMs can outperform token-level probabilities in some settings. We extend this line of work by measuring calibration in the context of a downstream escalation decision rather than confidence scores alone.

We draw on selective prediction (Geifman & El-Yaniv, 2017), and the related literature on learning to defer, where a classifier can route instances to a human expert (Madras et al., 2018; Mozannar & Sontag, 2020). Our cost-based formulation connects to cost-sensitive classification (Elkan, 2001), which optimizes decision thresholds under asymmetric misclassification costs. We adapt these ideas to LLM agents that must weigh confidence against explicit cost structures without access to class probabilities.

Chain-of-thought prompting (Wei et al., 2022) enables models to reason step-by-step, improving performance on arithmetic and multi-step tasks. Our finding that extended thinking is necessary for cost framing to take effect aligns with evidence that explicit reasoning unlocks capabilities latent in the base model. On the alignment side, RLHF (Ouyang et al., 2022) and DPO (Rafailov et al., 2023) train models to match human preferences; we use SFT with chain-of-thought supervision to align escalation behavior with the optimal policy.

Our multi-turn prompting protocol is related to multi-agent debate (Du et al., 2023; Liang et al., 2023), but the second turn evaluates the agent’s own prediction rather than introducing an adversarial perspective.

## 8 Conclusion

We have shown that LLM-based agents exhibit latent, model-specific escalation profiles that vary dramatically across models and cannot be predicted *a priori* from architecture or scale alone. Even within the same model family, scaling up produces unpredictable shifts: Gemma 3 4B and 12B differ by 22 percentage points in their implicit thresholds, and GPT-5-nano and GPT-5-mini differ by 38 points. Scaling dense Qwen3.5 from 4B to 9B barely moves the threshold but sharply worsens calibration and turns a step-like escalation curve into a noisy, gradual one. Models also carry miscalibrated beliefs about their own predictive accuracy, and the direction of miscalibration can reverse between a smaller and larger variant of the same family. Even calibration itself is model-specific: Claude Sonnet and Opus closely match their actual predictive accuracy, GPT-5-mini lands close to calibrated, while Gemma 3 models systematically underestimate it. Appendix F confirms these patterns extend to five additional models, including the Qwen3.5-397B flagship and the Llama and Mistral families. These latent decisions represent below-the-surface behavior that could disrupt a decision-making process without sufficient exploration.

These findings have immediate and extensive practical implications. Organizations deploying LLM agents for automated decision-making must empirically characterize their model’s

333 escalation behavior before deployment. Our methodology, based on varying predictive  
334 accuracy through calibrated signals, provides a tractable approach for doing so.

335 We also demonstrate that escalation behavior can be corrected through appropriate inter-  
336 ventions. The combination of extended thinking and cost framing achieves near-optimal  
337 escalation decisions, and supervised fine-tuning pushes escalation accuracy even higher  
338 while generalizing across domains, framings, and costs.

339 Several limitations suggest directions for future work. Our evaluation is limited to binary  
340 prediction tasks; real-world escalation decisions often involve more complex action spaces.  
341 We evaluate a relatively small number of models; a broader survey would strengthen the  
342 generality of our findings. Our cost structure assumes known, fixed costs; in practice, both  
343 error costs and human review costs may be uncertain or context-dependent.

344 **LLM Disclosure.** Large language models were used to help draft and edit this paper  
345 (including generating plots and figures) and to build the code repository for running  
346 experiments and analysis.

347 **Data and Code Availability.** All code, data, and experimental results are publicly available  
348 at <https://github.com/anonymous-code-submitter/colm-2026-escalation>.

## References

- Nuno Antonio, Ana de Almeida, and Luis Nunes. Hotel booking demand datasets. *Data in Brief*, 22:41–49, 2019. doi: 10.1016/j.dib.2018.11.126.
- Edmond Awad, Sohan Dsouza, Richard Kim, Jonathan Schulz, Joseph Henrich, Azim Shariff, Jean-François Bonnefon, and Iyad Rahwan. The moral machine experiment. *Nature*, 563: 59–64, 2018. doi: 10.1038/s41586-018-0637-6.
- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot arena: An open platform for evaluating LLMs by human preference. *arXiv preprint arXiv:2403.04132*, 2024.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.
- Charles Elkan. The foundations of cost-sensitive learning. In *International Joint Conference on Artificial Intelligence*, pp. 973–978, 2001.
- Yonatan Geifman and Ran El-Yaniv. Selective prediction in deep neural networks with design of experiments. In *Advances in Neural Information Processing Systems*, 2017.
- Nathan George. Lending club loan data. <https://www.kaggle.com/datasets/wordsforthewise/lending-club>, 2018.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning*, pp. 1321–1330, 2017.
- F. Maxwell Harper and Joseph A. Konstan. The MovieLens datasets: History and context. *ACM Transactions on Interactive Intelligent Systems*, 5(4):19:1–19:19, 2015. doi: 10.1145/2827872.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2024.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *arXiv preprint arXiv:2305.19118*, 2023.
- David Madras, Toniann Pitassi, and Richard Zemel. Predict responsibly: Improving fairness and accuracy by learning to defer. In *Advances in Neural Information Processing Systems*, 2018.
- Hussein Mozannar and David Sontag. Consistent estimators for learning to defer to an expert. In *International Conference on Machine Learning*, pp. 7076–7087, 2020.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2023.

- 397 Katherine Tian, Eric Mitchell, Huaxiu Yao, Christopher D. Manning, and Chelsea Finn. Just  
 398 ask for calibration: Strategies for eliciting calibrated confidence scores from language  
 399 models fine-tuned with human feedback. *arXiv preprint arXiv:2305.14975*, 2023.
- 400 Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin Manber, Vinber Dong, Edward  
 401 Li, Shashank Gupta, Ashish Sabharwal, and Niranjana Balasubramanian. AppWorld: A  
 402 controllable world of apps and people for benchmarking interactive coding agents. *arXiv*  
 403 *preprint arXiv:2407.18901*, 2024.
- 404 Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan  
 405 Chen, Jiakai Tang, Xu Chen, Yankai Lin, et al. A survey on large language model  
 406 based autonomous agents. *Frontiers of Computer Science*, 18(6), 2024. doi: 10.1007/  
 407 s11704-024-40231-1.
- 408 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi,  
 409 Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large  
 410 language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022.
- 411 Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang,  
 412 Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. AutoGen: Enabling next-gen LLM  
 413 applications via multi-agent conversation. In *COLM*, 2024.
- 414 Ellery Wulczyn, Nithum Thain, and Lucas Dixon. Ex machina: Personal attacks seen at  
 415 scale. In *Proceedings of the 26th International Conference on World Wide Web*, pp. 1391–1399,  
 416 2017. doi: 10.1145/3038912.3052591.
- 417 Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can  
 418 LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs.  
 419 *arXiv preprint arXiv:2306.13063*, 2024.

## 420 A *MoralMachine* results

- 421 *MoralMachine* presents ethical dilemmas rather than typical agent tasks, making it a robust-  
 422 ness check for whether escalation patterns generalize beyond standard decision-making.  
 423 GPT-5-nano and GPT-5-mini are excluded due to content filter restrictions.
- 424 Escalation rates are uniformly high across all models (68–100%), reflecting appropriate  
 425 uncertainty about ethical dilemmas where no objectively correct answer exists. The signal  
 426 has minimal effect: nohint escalation rates are similar to hint rates for most models, suggest-  
 427 ing that models recognize these decisions as inherently uncertain regardless of additional  
 428 information.

Table 2: *MoralMachine* escalation behavior. All models escalate at high rates, recognizing the inherent uncertainty of ethical dilemmas.

| Model             | With signal |           | Without signal |           |
|-------------------|-------------|-----------|----------------|-----------|
|                   | Pred. acc   | Esc. rate | Pred. acc      | Esc. rate |
| Qwen3.5-9B        | 61.2%       | 68.3%     | 52.1%          | 70.6%     |
| Qwen3.5-397B      | 70.4%       | 93.4%     | 69.3%          | 91.3%     |
| Llama 4 Maverick  | 65.8%       | 84.0%     | 65.2%          | 94.6%     |
| Llama 3.3 70B     | 67.5%       | 97.7%     | 71.4%          | 99.9%     |
| Mixtral 8x7B      | 65.1%       | 96.5%     | 39.5%          | 93.5%     |
| Mistral Small 24B | 70.2%       | 98.2%     | 68.4%          | 100.0%    |

## 429 B SFT signal ablation

- 430 To verify that the SFT model relies on the feature signal rather than exploiting some other  
 431 signal, we evaluate it on prompts that omit the signal entirely. Table 3 reports the results.

Table 3: SFT model evaluated with and without the signal, on the original cost framing.

| Dataset              | With signal | Without signal |
|----------------------|-------------|----------------|
| <i>HotelBookings</i> | 100%        | 84.7%          |
| <i>LendingClub</i>   | 100%        | 84.0%          |

Without the signal, escalation accuracy drops from 100% to approximately 84%. The model still follows the trained reasoning template, but it hallucinates a predictive-accuracy estimate (typically 75–85%) rather than reading it from the prompt. The hallucinated estimates are systematically optimistic, which produces correct escalation decisions at extreme cost ratios (where the decision is insensitive to the exact predictive accuracy) but incorrect decisions at  $R = 4$  (where escalation accuracy drops to 48–68%). The model’s arithmetic remains correct throughout; only the input is wrong. This confirms that the SFT model genuinely reads and uses the feature signal, and that the signal is the critical ingredient for perfect performance.

## C Proofs

*Proof of Theorem 1.* For a single instance with true accuracy  $p$ , the expected cost under threshold  $\tau$  is

$$C(p, \tau) = \begin{cases} c_\ell & \text{if } p < \tau \text{ (escalate)} \\ (1 - p) c_w & \text{if } p \geq \tau \text{ (implement).} \end{cases}$$

The agent should escalate when  $c_\ell < (1 - p) c_w$ , i.e., when  $p < 1 - c_\ell / c_w = \tau^*$ . This threshold uniquely minimizes expected cost because it is the only value at which the agent is indifferent between escalating and implementing.  $\square$

*Proof of Theorem 2.* The agent escalates when  $\hat{p} = p + \mu < \tau^*$ , i.e., when  $p < \tau^* - \mu$ . The agent therefore applies threshold  $\tau^* - \mu$ . The expected cost as a function of threshold  $\tau$  is  $C(\tau) = \int_0^\tau c_\ell f(p) dp + \int_\tau^1 (1 - p) c_w f(p) dp$ , with derivative  $C'(\tau) = f(\tau)[c_\ell - (1 - \tau)c_w]$ . For  $\tau < \tau^*$ , we have  $(1 - \tau) > c_\ell / c_w$ , so  $C'(\tau) < 0$ : moving the threshold further below  $\tau^*$  increases cost. For  $\tau > \tau^*$ ,  $C'(\tau) > 0$ : moving further above  $\tau^*$  also increases cost. Therefore  $C(\tau^* - \mu)$  is increasing in  $|\mu|$ .  $\square$

## D SFT results

## E Escalation threshold and self-estimated predictive accuracy

Figure 6 summarizes the implicit threshold  $p^*$  and self-estimated predictive accuracy  $\hat{a}$  for all thirteen models, as described in Sections 4 and 5.

## F Additional models

The main text reports eight models in four families (Qwen3.5, Gemma 3, Claude, GPT-5). We additionally evaluated five models using the same protocol: the Qwen3.5-397B-A17B MoE flagship, plus Llama 4 Maverick and Llama 3.3 70B (Meta), and Mixtral 8x7B and Mistral Small 24B (Mistral). The escalation-threshold and self-calibration patterns established in the main text extend to these models, with comparable within-family variation. Figure 6 reports the implicit threshold  $p^*$  and self-estimated predictive accuracy  $\hat{a}$  for all thirteen models side by side.

Qwen3.5-397B has  $p^* = 81\%$ , 25 percentage points above the dense Qwen3.5-9B ( $p^* = 56\%$ ), the largest within-Qwen threshold gap and a useful contrast to the near-flat 4B-to-9B threshold shift in the main text. Among the Llama and Mistral models,  $p^*$  exceeds 90% for all four, with Llama 3.3 70B and Mixtral 8x7B above 100%, and Mixtral 8x7B and



Table 4: SFT escalation accuracy for three fine-tuned models. Each is trained on three datasets and evaluated on all four main datasets, including the held-out *MovieLens*. Accuracy is measured against the Bayes-optimal policy across six cost ratios ( $R \in \{2, 4, 8, 10, 20, 50\}$ ), with 60 samples per (dataset, framing). Cells marked — were not evaluated (Qwen2.5-7B-Instruct predates the addition of the Study3 framing). *MoralMachine* results are reported in Appendix A.

| Model               | Dataset                          | Original | Dollar | Wording | Study3 |
|---------------------|----------------------------------|----------|--------|---------|--------|
| Qwen2.5-7B-Instruct | <i>HotelBookings</i> (train)     | 100%     | 98.3%  | 100%    | —      |
|                     | <i>LendingClub</i> (train)       | 100%     | 100%   | 100%    | —      |
|                     | <i>WikipediaToxicity</i> (train) | 100%     | 100%   | 100%    | —      |
|                     | <i>MovieLens</i> (held-out)      | 100%     | 100%   | 100%    | —      |
| Qwen3.5-4B          | <i>HotelBookings</i> (train)     | 98.3%    | 88.3%  | 98.3%   | 98.3%  |
|                     | <i>LendingClub</i> (train)       | 95.0%    | 95.0%  | 95.0%   | 90.0%  |
|                     | <i>WikipediaToxicity</i> (train) | 96.7%    | 98.3%  | 95.0%   | 96.7%  |
|                     | <i>MovieLens</i> (held-out)      | 98.3%    | 91.7%  | 98.3%   | 98.3%  |
| Qwen3.5-9B          | <i>HotelBookings</i> (train)     | 91.7%    | 98.3%  | 90.0%   | 88.3%  |
|                     | <i>LendingClub</i> (train)       | 93.3%    | 90.0%  | 86.7%   | 83.3%  |
|                     | <i>WikipediaToxicity</i> (train) | 76.7%    | 93.3%  | 73.3%   | 75.0%  |
|                     | <i>MovieLens</i> (held-out)      | 93.3%    | 85.0%  | 85.0%   | 93.3%  |




Figure 6: Implicit escalation threshold  $p^*$  (left) and self-estimated predictive accuracy  $\hat{a}$  (right) for each of the thirteen models. The threshold  $p^*$  spans 56% (Qwen3.5-9B) to nearly 100% (high-threshold models), while self-estimated predictive accuracy ranges from 68% to 97%. The dotted lines show the optimal threshold  $\tau^* = 75\%$  at cost ratio  $R = 4$  (left) and average actual accuracy (right).

Mistral Small 24B differ by more than 30 percentage points within the Mistral family. On the calibration axis, Mixtral 8x7B and Mistral Small 24B are strongly overconfident, Llama 3.3 70B is underconfident on aggregate but with offsetting per-condition biases, and Llama 4 Maverick falls in between. These results corroborate the main-text finding that escalation behavior is model-specific and unpredictable from architecture or scale.

## G Example scenarios

**HotelBookings.** “Person 1 has booked a hotel stay arriving on July 12, 2017 (week 28), with 2 weekend night(s) and 3 weekday night(s). The party consists of 2 adult(s). Person 1 is not a repeated guest and has 0 previous cancellation(s). They have requested 1 car parking space(s) and made 2 special request(s).”

478 **LendingClub.** “The applicant is requesting \$12,000 for debt consolidation. They have been  
479 employed for 5 years, a debt-to-income ratio of 14.3%, and a credit score of 712.”

480 **Wikipedia Toxicity.** “This comment needs to be checked: ‘Please stop removing content  
481 from Wikipedia. It is considered vandalism.’ ”

482 **MovieLens.** “Person 1 has reviewed: Toy Story (4/5), Braveheart (5/5), Apollo 13 (4/5),  
483 Pulp Fiction (5/5), Forrest Gump (3/5). Consider these two movies: Movie A: The Shaw-  
484 shank Redemption (Drama), average rating 4.43/5 (311 ratings). Movie B: Twelve Monkeys  
485 (Mystery—Sci-Fi—Thriller), average rating 3.64/5 (229 ratings).”

486 **MoralMachine.** “An autonomous vehicle is about to get in an accident. If the car doesn’t  
487 swerve, 2 elderly pedestrians will die. If the car swerves, 3 passengers will die. The  
488 pedestrians are legally crossing the street. The person behind the wheel is a 34-year-old  
489 male with a Bachelor Degree.”